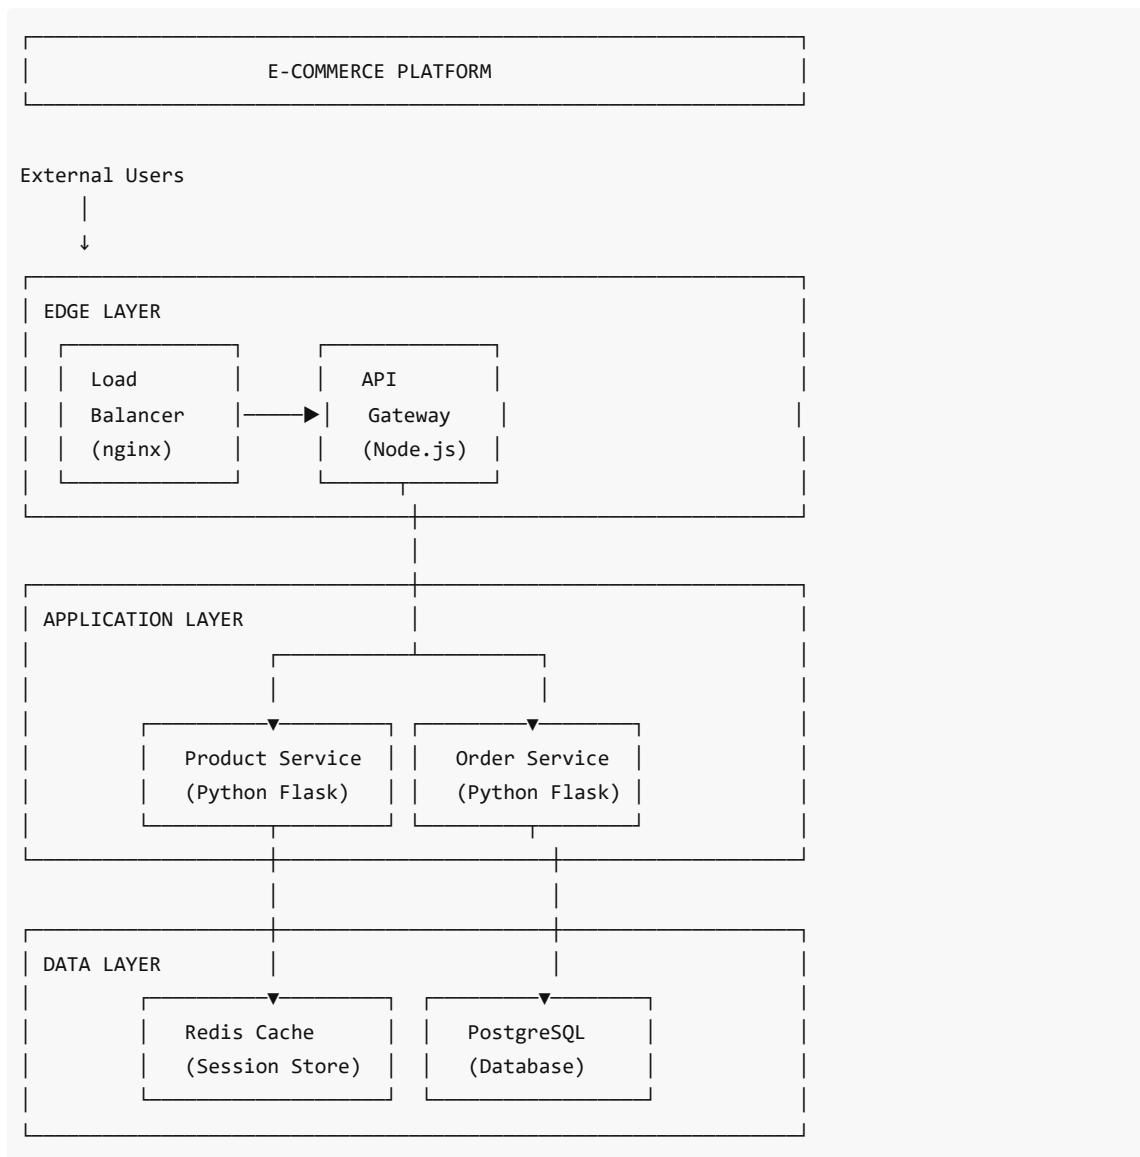


# Container Networking Project

## Project Overview

This project involves building a complete containerized microservices application infrastructure using only Linux primitives (network namespaces, veth pairs, bridges, iptables) to understand the low-level workings of container networking. The infrastructure simulates a real-world e-commerce platform with multiple services.

## System Architecture



## Prerequisites

- Linux environment (WSL2 or Linux VM)
- Root/Sudo privileges
- Python 3
- `iproute2` (for `ip` command)

- `iptables`
- `curl`

## 1: Foundation - Linux Primitives

### Goals

- Set up isolated network namespaces
- Create virtual network interfaces
- Implement basic inter-namespace communication

## Tasks

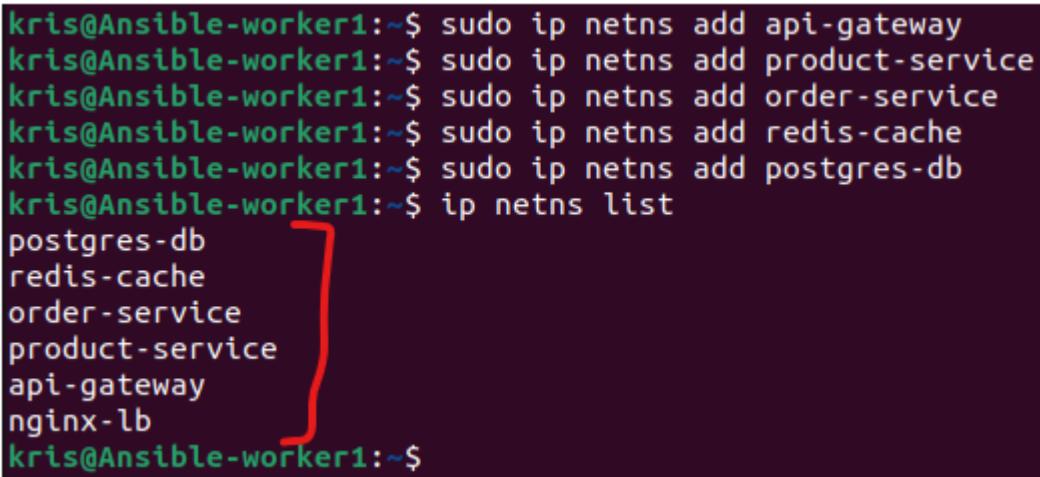
### Task 1.1: Create Network Namespaces

Create six network namespaces representing my services:

```
# Create namespaces
sudo ip netns add nginx-lb
sudo ip netns add api-gateway
sudo ip netns add product-service
sudo ip netns add order-service
sudo ip netns add redis-cache
sudo ip netns add postgres-db
```

### Deliverable: Screenshot showing all namespaces created

Run: `ip netns list`



```
kris@Ansible-worker1:~$ sudo ip netns add api-gateway
kris@Ansible-worker1:~$ sudo ip netns add product-service
kris@Ansible-worker1:~$ sudo ip netns add order-service
kris@Ansible-worker1:~$ sudo ip netns add redis-cache
kris@Ansible-worker1:~$ sudo ip netns add postgres-db
kris@Ansible-worker1:~$ ip netns list
postgres-db
redis-cache
order-service
product-service
api-gateway
nginx-lb
kris@Ansible-worker1:~$
```

### Task 1.2: Build a Virtual Bridge Network

---

Create a bridge to connect all services:

```
# Create bridge
sudo ip link add br-app type bridge
```

```
sudo ip addr add 10.0.0.1/16 dev br-app
sudo ip link set br-app up
```

**Connect each namespace to the bridge using veth pairs:**

```
# Example for nginx-lb (repeat for all services)
sudo ip link add veth-nginx type veth peer name veth-nginx-br
sudo ip link set veth-nginx netns nginx-lb
sudo ip link set veth-nginx-br master br-app
sudo ip link set veth-nginx-br up
```

**Configure inside namespace:**

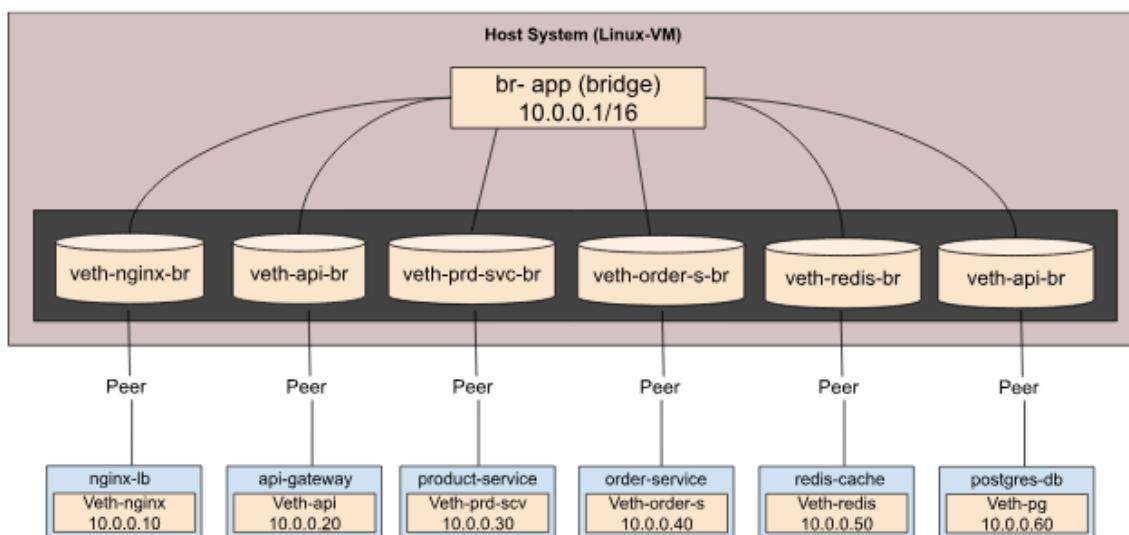
```
sudo ip netns exec nginx-lb ip addr add 10.0.0.10/16 dev veth-nginx
sudo ip netns exec nginx-lb ip link set veth-nginx up
sudo ip netns exec nginx-lb ip link set lo up
sudo ip netns exec nginx-lb ip route add default via 10.0.0.1
```

IP:

- nginx-lb: 10.0.0.10
- api-gateway: 10.0.0.20
- product-service: 10.0.0.30
- order-service: 10.0.0.40
- redis-cache: 10.0.0.50
- postgres-db: 10.0.0.60

**Deliverable:**

- Network diagram showing my setup



- Showing bridge with connected interfaces

```

kris@Ansible-worker1:~$ bridge link show
5: veth-nginx-br@if6: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 master br-app state forwarding priority 32 cost 2
7: veth-pg-br@if8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 master br-app state forwarding priority 32 cost 2
9: veth-redis-br@if10: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 master br-app state forwarding priority 32 cost 2
11: veth-order-s-br@if12: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 master br-app state forwarding priority 32 cost 2
13: veth-prd-svc-br@if14: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 master br-app state forwarding priority 32 cost 2
15: veth-api-br@if16: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 master br-app state forwarding priority 32 cost 2
kris@Ansible-worker1:~$ brctl show
bridge name     bridge id          STP enabled    interfaces
br-app          8000.faace6ed5435   no           veth-api-br
                                         veth-nginx-br
                                         veth-order-s-br
                                         veth-pg-br
                                         veth-prd-svc-br
                                         veth-redis-br

```

```

kris@Ansible-worker1:~$ ip link show type veth
5: veth-nginx-br@if6: <NOQUEUE,BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP mode DEFAULT group default qlen 1000
  link/ether 7a:0d:a1:08:94:d2 brd ff:ff:ff:ff:ff:ff link-metns nginx-lb
7: veth-pg-br@if8: <NOQUEUE,BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP mode DEFAULT group default qlen 1000
  link/ether 8a:b2:21:9c:34:17 brd ff:ff:ff:ff:ff:ff link-metns postgres-db
9: veth-redis-br@if10: <NOQUEUE,BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP mode DEFAULT group default qlen 1000
  link/ether e0:48:c1:be:bd:dc brd ff:ff:ff:ff:ff:ff link-metns redis-cache
11: veth-order-s-br@if12: <NOQUEUE,BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP mode DEFAULT group default qlen 1000
  link/ether 90:ea:db:f9:03:00 brd ff:ff:ff:ff:ff:ff link-metns order-service
13: veth-prd-svc-br@if14: <NOQUEUE,BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP mode DEFAULT group default qlen 1000
  link/ether 0a:dc:c4:a5:df:c0 brd ff:ff:ff:ff:ff:ff link-metns product-service
15: veth-api-br@if16: <NOQUEUE,BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP mode DEFAULT group default qlen 1000
  link/ether 5e:a3:ee:3d:27:13 brd ff:ff:ff:ff:ff:ff link-metns api-gateway
kris@Ansible-worker1:~$ 

```

- Proof of connectivity (ping tests between all namespaces)

Run: `sudo ip netns exec nginx-lb ping -c 2 10.0.0.20`

`nginx-lb` --> `api-gateway` (10.0.0.20)

`nginx-lb` --> `product-service`(10.0.0.30)

```

kris@Ansible-worker1:~$ sudo ip netns exec nginx-lb ping -c 2 10.0.0.20
PING 10.0.0.20 (10.0.0.20) 56(84) bytes of data.
64 bytes from 10.0.0.20: icmp_seq=1 ttl=64 time=0.043 ms
64 bytes from 10.0.0.20: icmp_seq=2 ttl=64 time=0.068 ms

--- 10.0.0.20 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1046ms
rtt min/avg/max/mdev = 0.043/0.055/0.068/0.012 ms
kris@Ansible-worker1:~$ sudo ip netns exec nginx-lb ping -c 2 10.0.0.30
PING 10.0.0.30 (10.0.0.30) 56(84) bytes of data.
64 bytes from 10.0.0.30: icmp_seq=1 ttl=64 time=0.971 ms
64 bytes from 10.0.0.30: icmp_seq=2 ttl=64 time=0.062 ms

--- 10.0.0.30 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 0.062/0.516/0.971/0.454 ms
kris@Ansible-worker1:~$ sudo ip netns exec nginx-lb ping -c 2 10.0.0.40
PING 10.0.0.40 (10.0.0.40) 56(84) bytes of data.
64 bytes from 10.0.0.40: icmp_seq=1 ttl=64 time=2.21 ms
64 bytes from 10.0.0.40: icmp_seq=2 ttl=64 time=0.080 ms

--- 10.0.0.40 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1002ms
rtt min/avg/max/mdev = 0.080/1.145/2.210/1.065 ms
kris@Ansible-worker1:~$ 

```

api-gateway --> product-service (10.0.0.30)

```
kris@Ansible-worker1:~$ sudo ip netns exec api-gateway ping -c 2 10.0.0.30
PING 10.0.0.30 (10.0.0.30) 56(84) bytes of data.
64 bytes from 10.0.0.30: icmp_seq=1 ttl=64 time=1.59 ms
64 bytes from 10.0.0.30: icmp_seq=2 ttl=64 time=0.053 ms

--- 10.0.0.30 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1002ms
rtt min/avg/max/mdev = 0.053/0.822/1.592/0.769 ms
kris@Ansible-worker1:~$ sudo ip netns exec api-gateway ping -c 2 10.0.0.40
PING 10.0.0.40 (10.0.0.40) 56(84) bytes of data.
64 bytes from 10.0.0.40: icmp_seq=1 ttl=64 time=1.34 ms
64 bytes from 10.0.0.40: icmp_seq=2 ttl=64 time=0.082 ms

--- 10.0.0.40 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1003ms
rtt min/avg/max/mdev = 0.082/0.711/1.341/0.629 ms
kris@Ansible-worker1:~$ █
```

### Task 1.3: Implement NAT for Internet Access

---

Enable internet access for all namespaces:

```
# Enable IP forwarding
sudo sysctl -w net.ipv4.ip_forward=1

# Add MASQUERADE rule
sudo iptables -t nat -A POSTROUTING -s 10.0.0.0/16 ! -o br-app -j MASQUERADE
```

### Deliverable: Test internet connectivity from each namespace

RUN: `sudo ip netns exec product-service ping -c 3 8.8.8.8`

```

kris@Ansible-worker1:~$ sudo ip netns exec product-service ping -c 3 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=254 time=76.1 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=254 time=79.5 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=254 time=82.6 ms

--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 76.056/79.378/82.571/2.661 ms
kris@Ansible-worker1:~$ sudo ip netns exec nginx-lb ping -c 3 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=254 time=83.3 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=254 time=74.6 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=254 time=78.8 ms

--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 74.599/78.897/83.334/3.567 ms
kris@Ansible-worker1:~$ sudo ip netns exec api-gateway ping -c 3 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=254 time=80.9 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=254 time=74.2 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=254 time=77.3 ms

--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 74.235/77.497/80.949/2.744 ms
kris@Ansible-worker1:~$ sudo ip netns exec order-service ping -c 3 8.8.8.8

```

```

kris@Ansible-worker1:~$ sudo ip netns exec order-service ping -c 3 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=254 time=88.8 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=254 time=76.5 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=254 time=80.6 ms

--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2001ms
rtt min/avg/max/mdev = 76.455/81.940/88.816/5.141 ms
kris@Ansible-worker1:~$ sudo ip netns exec redis-cache ping -c 3 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=254 time=78.8 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=254 time=79.7 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=254 time=73.2 ms

--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 73.165/77.218/79.711/2.891 ms
kris@Ansible-worker1:~$ sudo ip netns exec postgres-db ping -c 3 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=254 time=80.1 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=254 time=73.3 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=254 time=75.4 ms

--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 73.279/76.261/80.118/2.859 ms
kris@Ansible-worker1:~$ 

```

## Task 1.4: Setup Port Forwarding

**Forward host port 8080 to nginx-lb:**

```

# Add DNAT rule (PREROUTING - before routing decision)
sudo iptables -t nat -A PREROUTING -p tcp --dport 8080 -j DNAT --to-destination 10.0.0.10:80

# Allow forwarding in FILTER table
sudo iptables -A FORWARD -p tcp -d 10.0.0.10 --dport 80 -j ACCEPT

```

## Now external clients can access via host IP on port 8080

Test from host (simulating external client) to confirm request is forwarded to nginx-lb

RUN: curl http://localhost:8080

```
kris@Ansible-worker1:~$ curl http://localhost:8080
<!DOCTYPE html>
<html lang="en">
  <head>
    <meta charset="UTF-8" />
    <title>Apache Tomcat/8.5.53</title>
    <link href="favicon.ico" rel="icon" type="image/x-icon" />
    <link href="favicon.ico" rel="shortcut icon" type="image/x-icon" />
    <link href="tomcat.css" rel="stylesheet" type="text/css" />
  </head>
  <body>
    <div id="wrapper">
      <div id="navigation" class="curved container">
        <span id="nav-home"><a href="https://tomcat.apache.org/">Home</a></span>
        <span id="nav-hosts"><a href="/docs/">Documentation</a></span>
        <span id="nav-config"><a href="/docs/config/">Configuration</a></span>
        <span id="nav-examples"><a href="/examples/">Examples</a></span>
        <span id="nav-wiki"><a href="https://wiki.apache.org/tomcat/FrontPage">Wiki</a></span>
        <span id="nav-lists"><a href="https://tomcat.apache.org/lists.html">Mailing Lists</a></span>
        <span id="nav-help"><a href="https://tomcat.apache.org/findhelp.html">Find Help</a></span>
        <br class="separator" />
      </div>
      <div id="aspf-box">
```

## View the NAT rule:

RUN: sudo iptables -t nat -L -v -n

```
kris@Ansible-worker1:~$ sudo iptables -t nat -L -v -n
Chain PREROUTING (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target  prot opt in     out    source         destination
  0   0 DNAT    tcp   --  *      +      0.0.0.0/0          0.0.0.0/0          tcp dpt:8080 to:10.0.0.10:80

Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target  prot opt in     out    source         destination

Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target  prot opt in     out    source         destination

Chain POSTROUTING (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target  prot opt in     out    source         destination
 221 56394 MASQUERADE  all   --  *      !br-app  10.0.0.0/16    0.0.0.0/0

kris@Ansible-worker1:~$
```

## Deliverable: Document all iptables rules with explanations

### Explanation of the rule:

Rule:

1. sudo iptables -t nat -A POSTROUTING -s 10.0.0.0/16 ! -o br-app -j MASQUERADE

What this rule does:

- Allows containers/ namespaces to access the internet
- Hides internal IP addresses behind the host IP
- Required for outbound connectivity

2. sudo iptables -t nat -A PREROUTING -p tcp --dport 8080 -j DNAT --to-destination

```
10.0.0.10:80
```

What this rule does:

- Redirects traffic from host port 8080
- Sends it to nginx running inside the namespace on port 80
- Enables external access to internal services

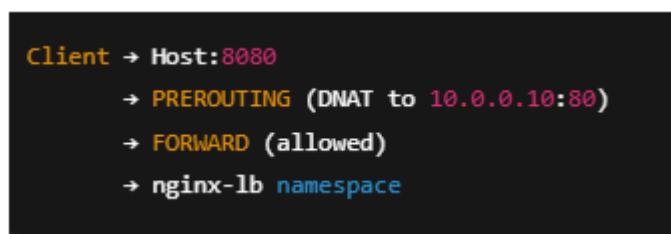
```
3. sudo iptables -A FORWARD -p tcp -d 10.0.0.10 --dport 80 -j ACCEPT
```

What this rule does:

- Explicitly allows forwarded traffic to reach nginx-lb
- Prevents Linux from dropping forwarded packets
- Required for DNAT port forwarding to function

#### Traffic Flow Summary:

External Client (Port 8080) -> Host (Port 8080) -> DNAT -> nginx-lb (Port 80)



Option / Flag	Meaning	Explanation
-t nat	NAT table	Specifies that the rule applies to the NAT table
-A POSTROUTING	POSTROUTING chain	Appends rule to POSTROUTING chain (applied <b>after</b> routing decision)
-A PREROUTING	PREROUTING chain	Appends rule to PREROUTING chain (applied <b>before</b> routing decision)
-A FORWARD	FORWARD chain	Appends rule to FORWARD chain (controls forwarded traffic)
-s 10.0.0.0/16	Source network	Matches packets originating from the container network
! -o br-app	Not bridge interface	Matches packets <b>not</b> leaving via the bridge interface (i.e., traffic going to the internet)
-p tcp	TCP protocol	Matches TCP traffic
--dport 8080	Destination port	Matches packets destined to port <b>8080</b> on the host
--dport 80	Destination port	Matches packets destined to port <b>80</b>
-d 10.0.0.10	Destination IP	Matches destination IP address of the nginx-lb namespace
-j MASQUERADE	SNAT action	Applies MASQUERADE (dynamic Source NAT), rewriting source IP to host's external IP

-j DNAT	DNAT action	Applies Destination NAT, rewriting destination IP and/or port
--to-destination 10.0.0.10:80	DNAT target	Fowards traffic to nginx-lb namespace at IP 10.0.0.10 on port 80
-j ACCEPT	Accept action	Allows the packet to be forwarded

## 2: Application Services

### Goals

- Deploy actual services in namespaces
- Implement service-to-service communication
- Test the complete application flow

### Tasks

#### Task 2.1: Deploy Nginx Load Balancer

Create a simple nginx configuration that load balances to the API gateway:

Install and run nginx in the namespace:

```
# Create nginx config

sudo ip netns exec nginx-lb bash -c 'cat <<EOF > /tmp/nginx/nginx.conf
events {
    worker_connections 1024;
}

http {
    upstream api_gateway {
        server 10.0.0.20:3000;
    }

    server {
        listen 80;

        location / {
            proxy_pass http://api_gateway;
            proxy_set_header Host \$host;
            proxy_set_header X-Real-IP \$remote_addr;
        }

        location /health {
            return 200 "OK\n";
            add_header Content-Type text/plain;
        }
    }
}
EOF'
```

Verify the nginx config file has been created in nginx-lb namespace

```
RUN: sudo ip netns exec nginx-lb cat /tmp/nginx/nginx.conf
```

Start nginx inside nginx-lb namespace

```
RUN: sudo ip netns exec nginx-lb nginx -c /tmp/nginx/nginx.conf
```

### Deliverable: Working load balancer responding to HTTP requests

**Test from host:**

```
RUN: curl http://192.168.56.104:8080/health
```

```
kris@Ansible-worker1:~$ curl http://192.168.56.104:8080/health
OK
kris@Ansible-worker1:~$
```

**Test from another namespace:**

```
RUN: sudo ip netns exec api-gateway curl http://10.0.0.10/health
```

```
kris@Ansible-worker1:~$ sudo ip netns exec api-gateway curl http://10.0.0.10/health
OK
kris@Ansible-worker1:~$ sudo ip netns exec redis-cache curl http://10.0.0.10/health
OK
kris@Ansible-worker1:~$ sudo ip netns exec postgres-db curl http://10.0.0.10/health
OK
kris@Ansible-worker1:~$ sudo ip netns exec product-service curl http://10.0.0.10/health
OK
kris@Ansible-worker1:~$ sudo ip netns exec order-service curl http://10.0.0.10/health
OK
kris@Ansible-worker1:~$
```

## Task 2.2: Create API Gateway

**Build a Node.js or Python API gateway that routes to backend services.**

**create api-gateway.py:**

```
from flask import Flask, jsonify, request
import requests

app = Flask(__name__)

PRODUCT_SERVICE = "http://10.0.0.30:5000"
ORDER_SERVICE = "http://10.0.0.40:5000"

@app.route('/health')
def health():
    return jsonify({"status": "healthy", "service": "api-gateway"})

@app.route('/api/products', methods=['GET'])
def get_products():
    try:
        response = requests.get(f"{PRODUCT_SERVICE}/products")
```

```

        return jsonify(response.json()), response.status_code
    except Exception as e:
        return jsonify({"error": str(e)}), 503

@app.route('/api/products/<id>', methods=['GET'])
def get_product(id):
    try:
        response = requests.get(f"{PRODUCT_SERVICE}/products/{id}")
        return jsonify(response.json()), response.status_code
    except Exception as e:
        return jsonify({"error": str(e)}), 503

@app.route('/api/orders', methods=['POST'])
def create_order():
    try:
        response = requests.post(
            f"{ORDER_SERVICE}/orders",
            json=request.json
        )
        return jsonify(response.json()), response.status_code
    except Exception as e:
        return jsonify({"error": str(e)}), 503

if __name__ == '__main__':
    app.run(host='0.0.0.0', port=3000)

```

Copy file to accessible location

```

sudo ip netns exec api-gateway bash -c 'cat <<EOF > /tmp/api-gateway.py
EOF'

# Install dependencies in namespace or use a Python virtual environment
sudo ip netns exec api-gateway pip install flask requests

# Start api-gateway
sudo ip netns exec api-gateway python3 /tmp/api-gateway.py &

```

#### **Flask running inside the api-gateway namespace on port 3000:**

RUN: sudo ip netns exec api-gateway ss -lntp | grep 3000

```

kris@Ansible-worker1:~$ sudo ip netns exec api-gateway ss -lntp | grep 3000
LISTEN 0      128          0.0.0.0:3000          0.0.0.0:*      users:(("python3",pid=314388,fd=3))
kris@Ansible-worker1:~$ 

```

#### **Testing the Api health from inside the namespace:**

RUN: sudo ip netns exec api-gateway curl http://127:3000/health

```
kris@Ansible-worker1:~$ sudo ip netns exec api-gateway curl http://127.0.0.1:3000/health
127.0.0.1 - - [26/Dec/2025 04:47:01] "GET /health HTTP/1.1" 200 -
{"service":"api-gateway","status":"healthy"}
kris@Ansible-worker1:~$
```

#### Test from other namespaces (e.g., nginx-lb)

RUN: `sudo ip netns exec nginx-lb curl http://10.0.0.20:3000/health`

```
kris@Ansible-worker1:~$ sudo ip netns exec nginx-lb curl http://10.0.0.20:3000/health
10.0.0.10 - - [26/Dec/2025 05:13:41] "GET /health HTTP/1.1" 200 -
{"service":"api-gateway","status":"healthy"}
kris@Ansible-worker1:~$
```

### Deliverable: API Gateway responding to requests and routing correctly

- Test API Gateway Routing to product-service from api-gateway namespace:

RUN: `sudo ip netns exec api-gateway curl http://localhost:3000/api/products`

```
kris@Ansible-worker1:~$ sudo ip netns exec api-gateway curl http://localhost:3000/api/products
[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]
kris@Ansible-worker1:~$
```

- Then test API Gateway Routing to product service via my host port-forwarding to confirm everything works end-to-end:

RUN: `curl http://192.168.56.104:8080/api/products`

```
kris@Ansible-worker1:~$ curl http://192.168.56.104:8080/api/products
[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]
kris@Ansible-worker1:~$
```

- Test API Gateway Routing to order-service from api-gateway namespace:

RUN: `sudo ip netns exec api-gateway curl http://localhost:3000/api/orders`

```
kris@Ansible-worker1:~$ sudo ip netns exec api-gateway curl http://localhost:3000/api/orders
<!DOCTYPE html>
<html lang=en>
<title>405 Method Not Allowed</title>
<h1>Method Not Allowed</h1>
<p>The method is not allowed for the requested URL.</p>
kris@Ansible-worker1:~$
```

- Then test API Gateway Routing to order service via my host port-forwarding to confirm everything works end-to-end:

RUN: `curl http://192.168.56.104:8080/api/orders`

```

kris@Ansible-worker1:~$ curl http://192.168.56.104:8080/api/orders
<!doctype html>
<html lang=en>
<title>405 Method Not Allowed</title>
<h1>Method Not Allowed</h1>
<p>The method is not allowed for the requested URL.</p>
kris@Ansible-worker1:~$ █

```

### Task 2.3: Build Product Service

Create product-service.py:

```

from flask import Flask, jsonify
import redis
import json

app = Flask(__name__)

# Connect to Redis cache
try:
    cache = redis.Redis(host='10.0.0.50', port=6379, decode_responses=True)
except:
    cache = None

# Mock product database
PRODUCTS = {
    "1": {"id": "1", "name": "Laptop", "price": 999.99, "stock": 50},
    "2": {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200},
    "3": {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150},
}

@app.route('/health')
def health():
    return jsonify({"status": "healthy", "service": "product-service"})

@app.route('/products', methods=['GET'])
def get_products():
    # Try cache first
    if cache:
        cached = cache.get('all_products')
        if cached:
            return jsonify(json.loads(cached))

    # Return products and cache
    products = list(PRODUCTS.values())
    if cache:
        cache.setex('all_products', 300, json.dumps(products))

    return jsonify(products)

@app.route('/products/<product_id>', methods=['GET'])

```

```

def get_product(product_id):
    # Try cache first
    if cache:
        cached = cache.get(f'product_{product_id}')
        if cached:
            return jsonify(json.loads(cached))

    product = PRODUCTS.get(product_id)
    if not product:
        return jsonify({"error": "Product not found"}), 404

    if cache:
        cache.setex(f'product_{product_id}', 300, json.dumps(product))

    return jsonify(product)

if __name__ == '__main__':
    app.run(host='0.0.0.0', port=5000)

```

**Start product-service:** sudo ip netns exec product-service python3 /tmp/product-service.py &

**Verify product-service is running:** sudo ip netns exec product-service ss -lntp | grep 5000

### Deliverable: Working product service with Redis caching

**Redis running inside its namespace and listening on port 6379:**

RUN: sudo ip netns exec redis-cache ss -lntp | grep 6379

```

kris@Ansible-worker1:~$ sudo ip netns exec redis-cache ss -lntp | grep 6379
LISTEN 0      511          0.0.0.0:6379      0.0.0.0:*      users:(("redis-server",pid=342464,fd=6))
kris@Ansible-worker1:~$ 

```

**Product-service successfully connect to redis cache:**

RUN: sudo ip netns exec product-service nc -zv 10.0.0.50 6379

```

kris@Ansible-worker1:~$ sudo ip netns exec product-service nc -zv 10.0.0.50 6379
Connection to 10.0.0.50 6379 port [tcp/redis] succeeded!
kris@Ansible-worker1:~$ 

```

**API Gateway Routing to product-service from api-gateway namespace:**

RUN: sudo ip netns exec api-gateway curl http://localhost:3000/api/products

```

kris@Ansible-worker1:~$ sudo ip netns exec api-gateway curl http://localhost:3000/api/products
[{"id": "1", "name": "Laptop", "price": 999.99, "stock": 50}, {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200}, {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150}]
kris@Ansible-worker1:~$ 

```

## Task 2.4: Build Order Service

Create `order-service.py`:

```
from flask import Flask, jsonify, request
import psycopg2
from datetime import datetime
import json

app = Flask(__name__)

# Database connection
def get_db():
    return psycopg2.connect(
        host='10.0.0.60',
        database='orders',
        user='postgres',
        password='postgres'
    )

# Initialize database
def init_db():
    conn = get_db()
    cur = conn.cursor()
    cur.execute('''
        CREATE TABLE IF NOT EXISTS orders (
            id SERIAL PRIMARY KEY,
            customer_id VARCHAR(100),
            product_id VARCHAR(100),
            quantity INTEGER,
            total_price DECIMAL(10, 2),
            created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
        )
    ''')
    conn.commit()
    cur.close()
    conn.close()

@app.route('/health')
def health():
    return jsonify({"status": "healthy", "service": "order-service"})

@app.route('/orders', methods=['POST'])
def create_order():
    data = request.json

    conn = get_db()
    cur = conn.cursor()

    cur.execute(
        '''INSERT INTO orders (customer_id, product_id, quantity, total_price)
           VALUES (%s, %s, %s, %s) RETURNING id''',
        (data['customer_id'], data['product_id'], data['quantity'], data['total_price']))
    result = cur.fetchone()
    conn.commit()
    cur.close()
    conn.close()

    return jsonify(result)
```

```

        (data['customer_id'], data['product_id'],
         data['quantity'], data['total_price'])
    )

    order_id = cur.fetchone()[0]
    conn.commit()
    cur.close()
    conn.close()

    return jsonify({"order_id": order_id, "status": "created"}), 201

@app.route('/orders/<order_id>', methods=['GET'])
def get_order(order_id):
    conn = get_db()
    cur = conn.cursor()

    cur.execute('SELECT * FROM orders WHERE id = %s', (order_id,))
    order = cur.fetchone()

    cur.close()
    conn.close()

    if not order:
        return jsonify({"error": "Order not found"}), 404

    return jsonify({
        "id": order[0],
        "customer_id": order[1],
        "product_id": order[2],
        "quantity": order[3],
        "total_price": float(order[4]),
        "created_at": order[5].isoformat()
    })
}

if __name__ == '__main__':
    init_db()
    app.run(host='0.0.0.0', port=5000)

```

**Start order-service:** sudo ip netns exec order-service python3 /tmp/order-service.py &

**Verify order-service is running and listening on port 5000:** sudo ip netns exec order-service ss -lntp | grep 5000

```
kris@Ansible-worker1:~$ sudo ip netns exec order-service ss -lntp | grep 5000
LISTEN 0      128      0.0.0.0:5000      0.0.0.0:*      users:(("python3",pid=341337,fd=3))
kris@Ansible-worker1:~$
```

**Deliverable: Working order service with PostgreSQL integration**

**postgres running as a container:** RUN: sudo docker ps

```
kris@Ansible-worker1:~$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS
3cb6c64af2c1        postgres:15        "docker-entrypoint.s..."   4 hours ago       Up 4 hours        8.0.0.0:5432->5432/tcp, [::]:5432->5432/tcp
kris@Ansible-worker1:~$
```

\*\*PostgreSQL listening on port 5432:\*\*  
RUN: `ss -lntp | grep 5432`

```
kris@Ansible-worker1:~$ ss -lntp | grep 5432
LISTEN 0      4096          0.0.0.0:5432          0.0.0.0:*
LISTEN 0      4096          [::]:5432            [::]:*
kris@Ansible-worker1:~$
```

#### Order-service reaching postgres:

RUN: `sudo ip netns exec order-service nc -zv 10.0.0.1 5432`

```
kris@Ansible-worker1:~$ sudo ip netns exec order-service nc -zv 10.0.0.1 5432
Connection to 10.0.0.1 5432 port [tcp/postgresql] succeeded!
kris@Ansible-worker1:~$
```

#### To confirm order-service can reach PostgreSQL and data is stored correctly:

- Creating an order (WRITE to PostgreSQL) to confirm order successfully created via API Gateway (PostgreSQL write)

```
sudo ip netns exec api-gateway curl -X POST http://localhost:3000/api/orders \
-H "Content-Type: application/json" \
-d '{
  "customer_id": "cust-101",
  "product_id": "prod-201",
  "quantity": 2,
  "total_price": 150.00
}'
```

```
kris@Ansible-worker1:~$ sudo ip netns exec api-gateway curl -X POST http://localhost:3000/api/orders \
-H "Content-Type: application/json" \
-d '{
  "customer_id": "cust-101",
  "product_id": "prod-201",
  "quantity": 2,
  "total_price": 150.00
}'
{"order_id":3,"status":"created"}
kris@Ansible-worker1:~$
```

- Verify order record exist and persisted in PostgreSQL database

RUN: `psql -U postgres -d orders -h 10.0.0.1 -c "SELECT * FROM orders;"`

```
kris@Ansible-worker1:~$ psql -U postgres -d orders -h 10.0.0.1 -c "SELECT * FROM orders;"  
Password for user postgres:  
id | customer_id | product_id | quantity | total_price | created_at  
---+-----+-----+-----+-----+-----+  
1 | cust-1 | prod-1 | 2 | 50.00 | 2025-12-27 10:33:52.142722  
2 | cust-1 | prod-1 | 2 | 50.00 | 2025-12-27 10:34:19.639925  
3 | cust-101 | prod-201 | 2 | 150.00 | 2025-12-27 12:08:19.329735  
(3 rows)  
  
kris@Ansible-worker1:~$
```

- check from within order-service namespace to confirm the order created via api gateway and stored in postgres:

RUN: sudo ip netns exec order-service psql -U postgres -d orders -h 10.0.0.1 -c "SELECT \* FROM orders;"

```
kris@Ansible-worker1:~$ sudo ip netns exec order-service psql -U postgres -d orders -h 10.0.0.1 -c "SELECT * FROM orders;"  
Password for user postgres:  
id | customer_id | product_id | quantity | total_price | created_at  
---+-----+-----+-----+-----+-----+  
1 | cust-1 | prod-1 | 2 | 50.00 | 2025-12-27 10:33:52.142722  
2 | cust-1 | prod-1 | 2 | 50.00 | 2025-12-27 10:34:19.639925  
3 | cust-101 | prod-201 | 2 | 150.00 | 2025-12-27 12:08:19.329735  
(3 rows)  
kris@Ansible-worker1:~$
```

## Task 2.5: Deploy Redis and PostgreSQL

**Run Redis:** sudo ip netns exec redis-cache redis-server --bind 0.0.0.0 &

```
kris@Ansible-worker1:~$ 345361:C 27 Dec 2025 04:29:51.242 # 000000000000 Redis is starting 000000000000  
345361:C 27 Dec 2025 04:29:51.242 # Redis version=6.0.16, bits=64, commit=00000000, modified=0, pid=345361, just started  
345361:C 27 Dec 2025 04:29:51.244 # Configuration loaded  
345361:M 27 Dec 2025 04:29:51.245 * Increased maximum number of open files to 10032 (it was originally set to 1024).  
345361:M 27 Dec 2025 04:29:51.247 # Could not create server TCP listening socket 0.0.0.0:6379: bind: Address already in use  
^C  
[1]+ Exit 1 sudo ip netns exec redis-cache redis-server --bind 0.0.0.0  
kris@Ansible-worker1:~$
```

**Run PostgreSQL:**

Create a PostgreSQL container on host:

```
docker run -d --name postgres \  
-p 5432:5432 \  
-e POSTGRES_USER=postgres \  
-e POSTGRES_PASSWORD=postgres \  
-e POSTGRES_DB=orders \  
postgres:15
```

RUN: sudo docker ps

```
kris@Ansible-worker1:~$ sudo docker ps  
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES  
3ch6c64af2c1 postgres:15 "docker-entrypoint.s...+" 4 hours ago Up 4 hours 0.0.0.0:5432->5432/tcp, [::]:5432->5432/tcp postgres  
kris@Ansible-worker1:~$
```

## Deliverable: Both data stores operational and accessible

- Redis running inside redis-cache namespace and accessible on port 6379

```
kris@Ansible-worker1:~$ sudo ip netns exec redis-cache ss -lntp | grep 6379
LISTEN 0      511          0.0.0.0:6379        0.0.0.0:*      users:(("redis-server",pid=342464,fd=6))
kris@Ansible-worker1:~$
```

- Product-service can reach Redis in redis-cache namespace

RUN: `sudo ip netns exec product-service nc -zv 10.0.0.50 6379`

```
kris@Ansible-worker1:~$ sudo ip netns exec product-service nc -zv 10.0.0.50 6379
Connection to 10.0.0.50 6379 port [tcp/redis] succeeded!
kris@Ansible-worker1:~$
```

- Order-service can reach postgresSQL

RUN: `sudo ip netns exec order-service nc -zv 10.0.0.1 5432`

```
kris@Ansible-worker1:~$ sudo ip netns exec order-service nc -zv 10.0.0.1 5432
Connection to 10.0.0.1 5432 port [tcp/postgresql] succeeded!
kris@Ansible-worker1:~$
```

## 3: Monitoring and Debugging

### Goals

- Implement network monitoring
- Create debugging tools
- Add observability to your infrastructure

### Tasks

#### Task 3.1: Network Traffic Analysis

Create a script that monitors traffic on the bridge and dump it on a log file:

```
#!/bin/bash
# monitor-traffic.sh

LOG_FILE="traffic_log_$(date '+%Y%m%d_%H%M%S').log"

echo "==== Network Traffic Monitor ==="
echo "Monitoring bridge: br-app"
echo "Logging to $LOG_FILE"
echo "Press Ctrl+C to stop"
```

```
echo ""

sudo tcpdump -i br-app -n -v -ttt > "$LOG_FILE"
```

**Run the traffic monitor:**

```
chmod +x monitor-traffic.sh

./monitor-traffic.sh
```

```
kris@Ansible-worker1:~$ ./monitor-traffic.sh
== Network Traffic Monitor ==
Monitoring bridge: br-app
Logging to traffic_log_20251227_053149.log
Press Ctrl+C to stop

tcpdump: listening on br-app, link-type EN10MB (Ethernet), snapshot length 262144 bytes
^C98 packets captured
98 packets received by filter
0 packets dropped by kernel
kris@Ansible-worker1:~$ ]
```

**Tcpdump running with app traffic:**

```
sudo tcpdump -i br-app -n -v -ttt
```

```
kris@Ansible-worker1:~$ sudo tcpdump -i br-app -n -v -ttt
(sudo) password for kris:
tcpdump: listening on br-app, link-type EN10MB (Ethernet), snapshot length 262144 bytes
2025-12-28 05:18:03.269312 IP (tos 0x0, ttl 64, id 39266, offset 0, flags [DF], proto TCP (6), length 60)
  10.0.0.1.46960 > 10.0.0.20.3800: Flags [S], cksm 0x1443 (incorrect -> 0xf3b6), seq 501080636, win 64240, options [mss 1468,sackOK,T5 val 4131687356 ecr 0,nop,wscale 7], length 0
2025-12-28 05:18:03.269435 IP (tos 0x0, ttl 64, id 0, offset 0, flags [DF], proto TCP (6), length 60)
  10.0.0.20.3800 > 10.0.0.1.46960: Flags [S.], cksm 0xi443 (incorrect -> 0x8462), seq 1334832442, ack 501080637, win 65160, options [mss 1468,sackOK,T5 val 10744851536 ecr 4131687356,nop,wscale 7], length 0
2025-12-28 05:18:03.269457 IP (tos 0x0, ttl 64, id 39267, offset 0, flags [DF], proto TCP (6), length 52)
  10.0.0.1.46960 > 10.0.0.20.3800: Flags [S.], cksm 0x143b (incorrect -> 0xaffc1), ack 1, win 582, options [nop,nop,T5 val 4131687356 ecr 1874851536], length 0
2025-12-28 05:18:03.292724 IP (tos 0x0, ttl 64, id 39268, offset 0, flags [DF], proto TCP (6), length 136)
  10.0.0.1.46960 > 10.0.0.20.3800: Flags [P.], cksm 0x148f (incorrect -> 0xi3a8), seq 1:85, ack 1, win 502, options [nop,nop,T5 val 4131687356 ecr 1874851536], length 0
```

**Create a python script to analyze my log file "traffic\_log" and create a csv file for each analysis**

```
sudo nano traffic_analysis_services.py
```

```
#!/usr/bin/env python3

import glob
import re
import pandas as pd
import matplotlib.pyplot as plt

# =====
# SERVICE ↔ IP MAP
# =====

SERVICE_MAP = {
    "10.0.0.10": "nginx-lb",
```

```

    "10.0.0.20": "api-gateway",
    "10.0.0.30": "product-service",
    "10.0.0.40": "order-service",
    "10.0.0.50": "redis-cache",
    "10.0.0.60": "postgres-db",
}

LOG_PATTERN = "traffic_log_*.log"

# =====
# LOAD & NORMALIZE TCPDUMP LOG
# =====

log_files = glob.glob(LOG_PATTERN)
if not log_files:
    print("✖ No traffic_log_*.log files found")
    exit(1)

raw_lines = []
for file in log_files:
    with open(file, "r") as f:
        raw_lines.extend(f.readlines())

print(f"✓ Loaded {len(raw_lines)} tcpdump lines")

# Merge indented lines with previous line
lines = []
buffer = ""

for line in raw_lines:
    if line.startswith(" ") or line.startswith("\t"):
        buffer += " " + line.strip()
    else:
        if buffer:
            lines.append(buffer)
        buffer = line.strip()

if buffer:
    lines.append(buffer)

# =====
# PARSE IPv4 PACKETS
# =====

IPV4_REGEX = re.compile(
    r"(\d+\.\d+\.\d+\.\d+)\.\d+\s*>\s*(\d+\.\d+\.\d+\.\d+)\.\d+"
)

records = []

for line in lines:
    if "IP6" in line:

```

```

        continue

    match = IPV4_REGEX.search(line)
    if not match:
        continue

    src_ip, dst_ip = match.groups()

    if "ICMP" in line:
        protocol = "ICMP"
    elif "UDP" in line:
        protocol = "UDP"
    elif "Flags" in line:
        protocol = "TCP"
    else:
        protocol = "OTHER"

    records.append({
        "source": SERVICE_MAP.get(src_ip, src_ip),
        "destination": SERVICE_MAP.get(dst_ip, dst_ip),
        "protocol": protocol
    })

# =====
# DATAFRAME
# =====

df = pd.DataFrame(records)

if df.empty:
    print("✗ No IPv4 service traffic detected")
    exit(1)

print(f"✓ Parsed {len(df)} packets")

# =====
# ANALYSIS
# =====

pair_counts = df.groupby(["source", "destination"]).size().reset_index(name="packets")
pair_counts.to_csv("packets_per_service_pair.csv", index=False)

protocol_counts = df["protocol"].value_counts().reset_index()
protocol_counts.columns = ["protocol", "packets"]
protocol_counts.to_csv("protocol_distribution.csv", index=False)

top_talkers = df["source"].value_counts().reset_index()
top_talkers.columns = ["service", "packets"]
top_talkers.to_csv("top_talkers.csv", index=False)

# =====
# GRAPHS

```

```

# =====

plt.figure()
plt.bar(protocol_counts["protocol"], protocol_counts["packets"])
plt.title("Protocol Distribution")
plt.tight_layout()
plt.savefig("protocol_distribution.png")
plt.close()

top_pairs = pair_counts.sort_values("packets", ascending=False).head(10)
plt.figure(figsize=(10, 5))
plt.barh(
    top_pairs["source"] + " → " + top_pairs["destination"],
    top_pairs["packets"]
)
plt.title("Top Service-to-Service Traffic")
plt.tight_layout()
plt.savefig("top_service_pairs.png")
plt.close()

plt.figure()
plt.bar(top_talkers["service"], top_talkers["packets"])
plt.title("Top Talkers")
plt.tight_layout()
plt.savefig("top_talkers.png")
plt.close()

print("\n🎉 Traffic analysis complete!")

```

**Run python file: `python3 traffic_analysis_services.py` to analyze the tcldump log file and create csv file for each analysis**

```

kris@Ansible-worker1:~$ python3 traffic_analysis_services.py
✓ Loaded 234 tcldump lines
✓ Parsed 86 packets

🎉 Traffic analysis complete!
kris@Ansible-worker1:~$ █

```

**CSV files created by python script:**

```

kris@Ansible-worker1:~$ ls | grep .csv
packets_per_service_pair.csv
protocol_distribution.csv
top_talkers.csv
traffic_summary.csv
kris@Ansible-worker1:~$ █

```

**Assignment: Create a traffic analysis report showing:**

- Packets between each service pair

```
kris@Ansible-worker1:~$ cat packets_per_service_pair.csv
source,destination,packets
192.168.56.104,nginx-lb,18
api-gateway,nginx-lb,15
api-gateway,product-service,7
nginx-lb,192.168.56.104,12
nginx-lb,api-gateway,21
product-service,api-gateway,5
product-service,redis-cache,5
redis-cache,product-service,3
kris@Ansible-worker1:~$ ]
```

- Protocol distribution (TCP/UDP/ICMP)

```
kris@Ansible-worker1:~$ cat protocol_distribution.csv
protocol,packets
TCP,86
kris@Ansible-worker1:~$ ]
```

- Top talkers (most active services)

```
kris@Ansible-worker1:~$ cat top_talkers.csv
service,packets
nginx-lb,33
api-gateway,22
192.168.56.104,18
product-service,10
redis-cache,3
kris@Ansible-worker1:~$ ]
```

**Deliverable: Traffic analysis report with graphs:**

Create a python script `sudo nano generate_traffic_report.py` file that will generate a traffic analysis report for each csv file created by the python script **traffic\_analysis\_services.py**.

```
#!/usr/bin/env python3

from reportlab.platypus import (
    SimpleDocTemplate, Paragraph, Spacer, Image, Table, TableStyle
)
from reportlab.lib.styles import getSampleStyleSheet
from reportlab.lib.pagesizes import A4
from reportlab.lib import colors
```

```

import pandas as pd
import os
from datetime import datetime

# =====
# FILES
# =====

REPORT_NAME = "traffic_analysis_report.pdf"

CSV_FILES = {
    "Service Traffic": "packets_per_service_pair.csv",
    "Protocol Distribution": "protocol_distribution.csv",
    "Top Talkers": "top_talkers.csv",
}

IMAGES = [
    "protocol_distribution.png",
    "top_service_pairs.png",
    "top_talkers.png",
]

# =====
# DOCUMENT SETUP
# =====

doc = SimpleDocTemplate(REPORT_NAME, pagesize=A4)
styles = getSampleStyleSheet()
elements = []

# =====
# TITLE
# =====

elements.append(Paragraph(
    "<b>Network Traffic Analysis Report</b>",
    styles["Title"]
))
elements.append(Spacer(1, 12))

elements.append(Paragraph(
    f"Generated on: {datetime.now().strftime('%Y-%m-%d %H:%M:%S')}",
    styles["Normal"]
))
elements.append(Spacer(1, 20))

# =====
# EXECUTIVE SUMMARY
# =====

summary_text = """
This report presents an analysis of network traffic captured on the <b>br-app</b>

```

```

Linux bridge connecting multiple microservices implemented using network namespaces.
The goal is to validate service-to-service communication, protocol usage, and identify
top traffic sources within the application architecture.

"""

elements.append(Paragraph("<b>Executive Summary</b>", styles["Heading2"]))
elements.append(Spacer(1, 8))
elements.append(Paragraph(summary_text, styles["Normal"]))
elements.append(Spacer(1, 20))

# =====
# TABLE SECTIONS
# =====

for title, csv_file in CSV_FILES.items():
    if not os.path.exists(csv_file):
        continue

    df = pd.read_csv(csv_file)

    elements.append(Paragraph(f"<b>{title}</b>", styles["Heading2"]))
    elements.append(Spacer(1, 10))

    table_data = [df.columns.tolist()] + df.values.tolist()

    table = Table(table_data, repeatRows=1)
    table.setStyle(TableStyle([
        ("BACKGROUND", (0, 0), (-1, 0), colors.lightgrey),
        ("GRID", (0, 0), (-1, -1), 0.5, colors.grey),
        ("ALIGN", (0, 0), (-1, -1), "CENTER"),
        ("FONT", (0, 0), (-1, 0), "Helvetica-Bold"),
    ]))

    elements.append(table)
    elements.append(Spacer(1, 20))

# =====
# GRAPHS
# =====

elements.append(Paragraph("<b>Traffic Visualization</b>", styles["Heading2"]))
elements.append(Spacer(1, 12))

for img in IMAGES:
    if os.path.exists(img):
        elements.append(Image(img, width=400, height=250))
        elements.append(Spacer(1, 15))

# =====
# BUILD REPORT
# =====

```

```
doc.build(elements)

print("✅ PDF report generated successfully:")
print(f"  {REPORT_NAME}")
```

## Network Traffic Analysis Report

Generated on: 2025-12-28 06:23:18

### Executive Summary

This report presents an analysis of network traffic captured on the **br-app** Linux bridge connecting multiple microservices implemented using network namespaces. The goal is to validate service-to-service communication, protocol usage, and identify top traffic sources within the application architecture.

### Service Traffic

source	destination	packets
192.168.56.104	nginx-lb	18
api-gateway	nginx-lb	15
api-gateway	product-service	7
nginx-lb	192.168.56.104	12
nginx-lb	api-gateway	21
product-service	api-gateway	5
product-service	redis-cache	5
redis-cache	product-service	3

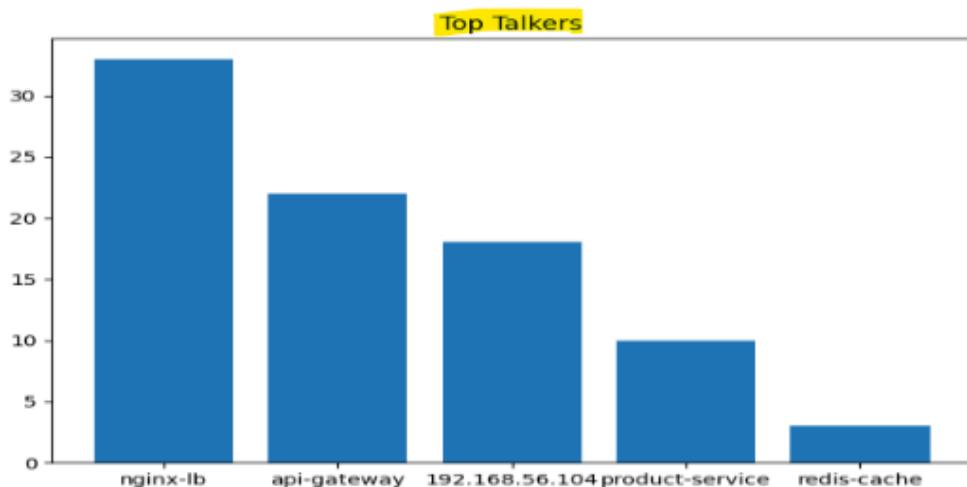
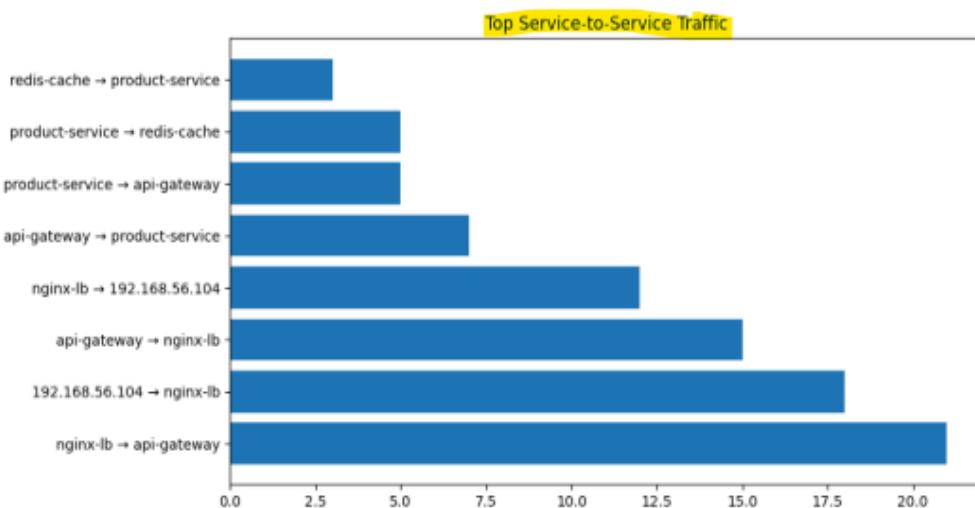
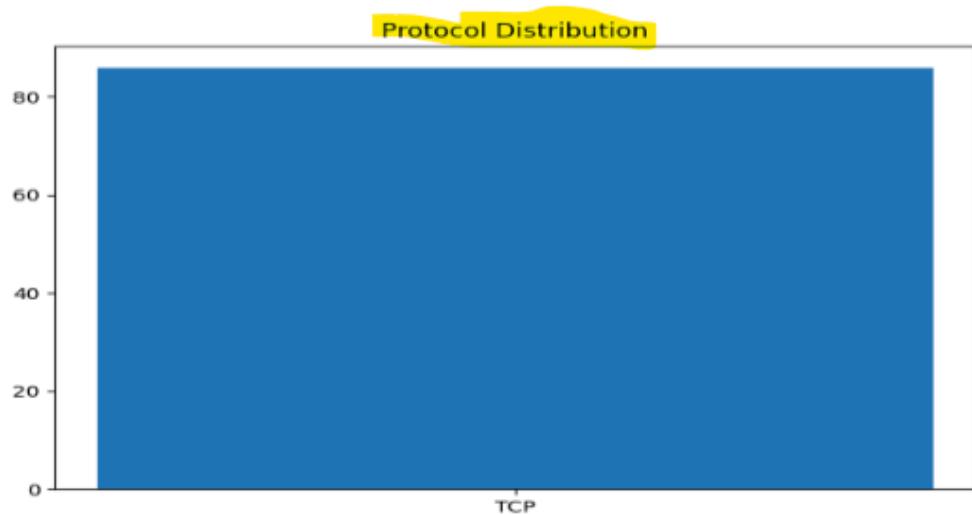
### Protocol Distribution

protocol	packets
TCP	86

### Top Talkers

service	packets
nginx-lb	33
api-gateway	22
192.168.56.104	18
product-service	10
redis-cache	3

## Traffic Visualization



## Task 3.2: Service Health Monitoring

Create health-monitor.py:

```
#!/usr/bin/env python3
import requests
import time
from datetime import datetime
import json

SERVICES = {
    'nginx-lb': 'http://10.0.0.10/health',
    'api-gateway': 'http://10.0.0.20:3000/health',
    'product-service': 'http://10.0.0.30:5000/health',
    'order-service': 'http://10.0.0.40:5000/health',
}

def check_health(service_name, url):
    try:
        response = requests.get(url, timeout=2)
        if response.status_code == 200:
            return {"status": "UP", "latency": response.elapsed.total_seconds()}
        else:
            return {"status": "DOWN", "error": f"HTTP {response.status_code}"}
    except Exception as e:
        return {"status": "DOWN", "error": str(e)}

def monitor():
    print("== Service Health Monitor ==")
    print(f"Started at: {datetime.now()}")
    print("")

    while True:
        print(f"\n[{datetime.now().strftime('%H:%M:%S')}] Health Check:")
        print("-" * 60)

        for service, url in SERVICES.items():
            health = check_health(service, url)
            status_symbol = "✓" if health['status'] == 'UP' else "✗"

            if health['status'] == 'UP':
                print(f"{status_symbol} {service:20s} UP  (latency: {health['latency']*1000:.2f}ms)")
            else:
                print(f"{status_symbol} {service:20s} DOWN ({health.get('error', 'Unknown')})")

        time.sleep(10)

if __name__ == '__main__':
    monitor()
```

## Deliverable: Health monitoring dashboard showing service status

Modify file permission and execute:

```
sudo chmod +x health-monitor.py
```

```
sudo ./health-monitor.py
```

```
kris@Ansible-worker1:~/git/ansible-tower-project$ ./health-monitor.py
== Service Health Monitor ==
Started at: 2025-12-28 07:21:17.836283

[07:21:17] Health Check:
-----
✓ nginx-lb           UP   (latency: 2.62ms)
✓ api-gateway        UP   (latency: 17.28ms)
✓ product-service    UP   (latency: 31.55ms)
✓ order-service      UP   (latency: 16.98ms)

[07:21:27] Health Check:
-----
✓ nginx-lb           UP   (latency: 3.54ms)
✓ api-gateway        UP   (latency: 3.63ms)
✓ product-service    UP   (latency: 35.83ms)
✓ order-service      UP   (latency: 3.15ms)

[07:21:38] Health Check:
-----
✓ nginx-lb           UP   (latency: 135.64ms)
✓ api-gateway        UP   (latency: 62.47ms)
✓ product-service    UP   (latency: 30.66ms)
✓ order-service      UP   (latency: 11.18ms)
```

## Task 3.3: Connection Tracking Analysis

Create a script to analyze active connections: `sudo nano connection-tracker.sh`

```
#!/bin/bash
# connection-tracker.sh

echo "== Active Connection Tracker =="
echo "Press Ctrl+C to stop"

while true; do
    clear
    echo "== Active Connections $(date) =="
    echo ""

    echo "Connections by Service:"
    echo "-----"

    for ns in nginx-lb api-gateway product-service order-service; do
        if ip netns list | grep -q "$ns"; then
            count=$(sudo ip netns exec $ns ss -tan state established | wc -l)
            echo "$ns: $count active connections"
        fi
    done
done
```

```

else
    echo "$ns: namespace not found"
fi
done

echo ""
echo "Connection States (conntrack):"
echo "-----"

sudo conntrack -L 2>/dev/null | grep "10.0.0" | \
awk '
{
    for(i=1;i<=NF;i++) {
        if($i ~ /^dst=/) {
            gsub("dst=", "", $i)
            print $i
        }
    }
}' | \
sed \
-e 's/10.0.0.10/nginx-lb/' \
-e 's/10.0.0.20/api-gateway/' \
-e 's/10.0.0.30/product-service/' \
-e 's/10.0.0.40/order-service/' \
-e 's/10.0.0.50/redis-cache/' | \
sort | uniq -c | sort -rn

sleep 5
done

```

## **Deliverable: Connection tracking report**

---

Modify file permission and execute

```

chmod +x connection-tracker.sh

sudo ./connection-tracker.sh

```

```
== Active Connections (Sun Dec 28 07:58:12 AM PST 2025) ==

Connections by Service:
-----
nginx-lb: 1 active connections
api-gateway: 1 active connections
product-service: 2 active connections
order-service: 1 active connections

Connection States (conntrack):
-----
 201 10.0.0.1
 68 product-service
 67 order-service
 67 api-gateway
 1 redis-cache
^C
```

# Connection Tracking Analysis Report

Date: 2025-12-28 16:17:09

Task: Task 3.3 – Connection Tracking Analysis

## Overview

This report presents an analysis of active TCP connections within the containerized microservices environment using Linux network namespaces and conntrack. The objective is to validate inter-service connectivity, traffic distribution, and backend dependency usage.

## Active Connections by Service

Service	Active Connections
nginx-lb	1
api-gateway	1
product-service	2
order-service	1

## Connection Distribution (conntrack)

Destination	Connection Count
Host (10.0.0.1)	201
product-service	68
order-service	67
api-gateway	67
redis-cache	1

## Analysis & Observations

- API Gateway and backend services show balanced and stable connection usage.
- Product and Order services maintain consistent active connections, validating service-to-service calls.
- Redis shows minimal connections, consistent with cache-access patterns.
- Host-level traffic confirms ingress routing through the bridge network.

## Conclusion

The connection tracking results confirm a healthy microservices networking setup. All services are reachable, connections are stable, and conntrack correctly records active flows across namespaces. This satisfies the deliverable requirements for Task 3.3.

## Task 3.4: Network Topology Visualizer

Create a script that generates a visual representation of your network: `sudo nano topology-visualizer.py`

```
#!/usr/bin/env python3
# topology-visualizer.py

import subprocess
```

```

import re

NAMESPACES = [
    'nginx-lb',
    'api-gateway',
    'product-service',
    'order-service',
    'redis-cache',
    'postgres-db'
]

def get_namespace_ips():
    """Get IPv4 addresses for all namespaces"""
    ips = {}

    for ns in NAMESPACES:
        try:
            result = subprocess.run(
                ['sudo', 'ip', 'netns', 'exec', ns, 'ip', '-4', 'addr', 'show'],
                capture_output=True, text=True, check=True
            )
            match = re.search(r'inet (\d+\.\d+\.\d+\.\d+)', result.stdout)
            if match:
                ips[ns] = match.group(1)
            else:
                ips[ns] = "N/A"
        except subprocess.CalledProcessError:
            ips[ns] = "N/A"

    return ips

def draw_topology():
    ips = get_namespace_ips()

    print("=" * 80)
    print(" " * 25 + "NETWORK TOPOLOGY DIAGRAM")
    print("=" * 80)
    print()
    print("          ● Internet")
    print("          |")
    print("          | NAT")
    print("          ▼")
    print("          |")
    print("          | Host")
    print("          | Port: 8080")
    print("          |")
    print("          | DNAT")
    print("          ▼")
    print("          |")
    print("          | Bridge: br-app")
    print("          | IP: 10.0.0.1")
    print("          |")

```

```
print("                                     | ")
print("                                     | ")
print("                                     | ")           | "
for service, ip in ips.items():
    print(f" {service:18s} → {ip}")
print()
print("=" * 80)
if __name__ == "__main__":
    draw_topology()
```

## **Deliverable: Network topology diagram**

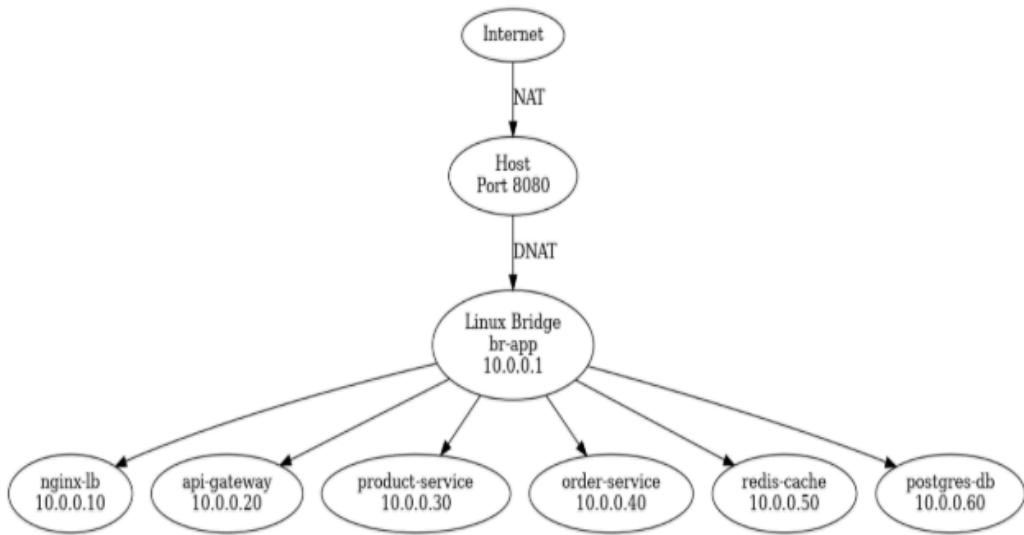
Modify file permission and execute:

```
sudo chmod +x topology-visualizer.py
```

```
python3 topology-visualizer.py
```

```
kris@Ansible-worker1:~$ nano topology-visualizer.py
kris@Ansible-worker1:~$ chmod +x topology-visualizer.py
kris@Ansible-worker1:~$ python3 topology-visualizer.py
=====
                    NETWORK TOPOLOGY
=====

    Internet
        |
        ↓ (NAT)
        |
    Host
    Port: 8080
        |
        ↓ (DNAT)
        |
Bridge: br-app
IP: 10.0.0.1
        |
        |-----|
nginx-lb      127.0.0.1
api-gateway   127.0.0.1
product-service 127.0.0.1
order-service   127.0.0.1
redis-cache     127.0.0.1
postgres-db     127.0.0.1
=====
```



## 4: Advanced Networking

### Goals

- Implement service discovery
- Add load balancing
- Create network security policies

### Tasks

#### Task 4.1: Implement Simple Service Discovery

Create a simple DNS-like service registry: `sudo nano service-registry.py`

```

#!/usr/bin/env python3
# service-registry.py

from flask import Flask, jsonify, request
import time

app = Flask(__name__)

# In-memory registry
services = {}

@app.route('/register', methods=['POST'])
def register_service():
    data = request.json

    service_name = data['name']
    service_ip = data['ip']
    service_port = data['port']

    services[service_name] = {
        'ip': service_ip,
        'port': service_port
    }

```

```

services[service_name] = {
    "ip": service_ip,
    "port": service_port,
    "registered_at": time.strftime('%Y-%m-%d %H:%M:%S'),
    "health": "unknown"
}

return jsonify({
    "status": "registered",
    "service": service_name
})

@app.route('/discover/<service_name>', methods=['GET'])
def discover_service(service_name):
    if service_name in services:
        return jsonify(services[service_name])
    return jsonify({"error": "Service not found"}), 404

@app.route('/services', methods=['GET'])
def list_services():
    return jsonify(services)

@app.route('/deregister/<service_name>', methods=['DELETE'])
def deregister(service_name):
    if service_name in services:
        del services[service_name]
        return jsonify({"status": "deregistered"})
    return jsonify({"error": "Service not found"}), 404

if __name__ == "__main__":
    app.run(host="10.0.0.70", port=8500)

```

### Run registry in its own namespace

Create a new namespace and start registry:

```
sudo ip netns exec service-registry python3 service-registry.py
```

```

kris@Ansible-worker1: ~ $ sudo ip netns add service-registry
sudo ip link add veth-sr type veth peer name veth-sr-br
sudo ip link set veth-sr netns service-registry
sudo ip link set veth-sr-br master br-app
sudo ip link set veth-sr-br up

sudo ip netns exec service-registry ip addr add 10.0.0.70/24 dev veth-sr
sudo ip netns exec service-registry ip link set veth-sr up
sudo ip netns exec service-registry ip link set lo up
kris@Ansible-worker1: ~ $ sudo ip netns exec service-registry python3 service-registry.py
* Serving Flask app 'service-registry'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:8500
* Running on http://127.0.0.1:8500
Press CTRL+C to quit

```

### Edit my services file (product, order, api-gateway) to register themselves on startup

```
sudo ip netns exec product-service python3 product-service.py
```

```
sudo ip netns exec order-service python3 order-service.py
```

```
sudo ip netns exec api-gateway python3 api-gateway.py
```

```
kris@Ansible-worker1: $ sudo ip netns exec service-registry curl http://10.0.0.70:8500/services
[sudo] password for kris:
()
kris@Ansible-worker1: $ sudo ip netns exec product-service python3 product-service.py
✓ Registered product-service with service registry
* Serving Flask app 'product-service'
* Debug mode: off
Address already in use
Port 5000 is in use by another program. Either identify and stop that program, or start the server with a different port.
kris@Ansible-worker1: $ sudo ip netns exec api-gateway python3 api-gateway.py
✓ Registered api-gateway with service registry
* Serving Flask app 'api-gateway'
* Debug mode: off
Address already in use
Port 3000 is in use by another program. Either identify and stop that program, or start the server with a different port.
kris@Ansible-worker1: $ sudo ip netns exec order-service python3 order-service.py
✓ Registered order-service with service registry
* Serving Flask app 'order-service'
* Debug mode: off
Address already in use
Port 5000 is in use by another program. Either identify and stop that program, or start the server with a different port.
kris@Ansible-worker1: $
```

## Deliverable: Working service discovery with all services registered

```
sudo ip netns exec service-registry python3 service-registry.py
```

```
kris@Ansible-worker1: $ sudo ip netns exec service-registry python3 service-registry.py
* Serving Flask app 'service-registry'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on http://10.0.0.70:8500
Press CTRL+C to quit
10.0.0.70 - - [29/Dec/2025 14:31:16] "GET /services HTTP/1.1" 200 -
10.0.0.1 - - [29/Dec/2025 14:35:16] "GET /services HTTP/1.1" 200 -
10.0.0.1 - - [29/Dec/2025 14:35:46] "GET /discover/product-service HTTP/1.1" 404 -
10.0.0.1 - - [29/Dec/2025 14:35:57] "GET /services HTTP/1.1" 200 -
10.0.0.40 - - [29/Dec/2025 14:37:05] "POST /register HTTP/1.1" 200 -
10.0.0.30 - - [29/Dec/2025 14:38:37] "POST /register HTTP/1.1" 200 -
10.0.0.20 - - [29/Dec/2025 14:39:21] "POST /register HTTP/1.1" 200 - ]
```

## Test Service Discovery

Run: curl http://10.0.0.70:8500/services

```
kris@Ansible-worker1: $ curl http://10.0.0.70:8500/services
{"api-gateway":{"health":"unknown","ip":"10.0.0.20","port":3000,"registered_at":"2025-12-29 14:39:21"}, "order-service":{"health":"unknown","ip":"10.0.0.40","port":5000,"registered_at":"2025-12-29 14:41:20"}, "product-service":{"health":"unknown","ip":"10.0.0.30","port":5000,"registered_at":"2025-12-29 14:38:37"}}
kris@Ansible-worker1: $
```

Run curl http://10.0.0.70:8500/discover/product-service

Run curl http://10.0.0.70:8500/discover/order-service

Run curl http://10.0.0.70:8500/discover/api-gateway

```

kris@Ansible-worker1: $ curl http://10.0.0.70:8500/discover/product-service
{"health": "unknown", "ip": "10.0.0.30", "port": 5000, "registered_at": "2025-12-29 14:38:37"}
kris@Ansible-worker1: $ curl http://10.0.0.70:8500/discover/order-service
{"health": "unknown", "ip": "10.0.0.40", "port": 5000, "registered_at": "2025-12-29 14:41:20"}
kris@Ansible-worker1: $ curl http://10.0.0.70:8500/discover/api-gateway
{"health": "unknown", "ip": "10.0.0.20", "port": 3000, "registered_at": "2025-12-29 14:39:21"}
kris@Ansible-worker1:-

```

## Task 4.2: Implement Round-Robin Load Balancing

**Modify the API Gateway to load balance between multiple instances:**

```

# Enhanced API Gateway with load balancing
import requests
import itertools

class LoadBalancer:
    def __init__(self, backends):
        self.backends = itertools.cycle(backends)
        self.backend_list = backends

    def get_backend(self):
        return next(self.backends)

    def health_check(self):
        healthy = []
        for backend in self.backend_list:
            try:
                response = requests.get(f"{backend}/health", timeout=1)
                if response.status_code == 200:
                    healthy.append(backend)
            except:
                pass
        self.backends = itertools.cycle(healthy)
        return healthy

# Use in routes
product_lb = LoadBalancer([
    "http://10.0.0.30:5000",
    "http://10.0.0.31:5000", # Add second instance
    "http://10.0.0.32:5000", # Add third instance
])

```

Create multiple instances of product-service

```

#!/bin/bash
# multiple_instance.sh
# Adds extra network namespaces for load balancing testing

set -e

```

```

echo "Setting up extra product service nodes..."

# Create namespaces
sudo ip netns add product2 || true
sudo ip netns add product3 || true

# Create veth pairs
sudo ip link add veth-prod2 type veth peer name veth-prod2-br || true
sudo ip link add veth-prod3 type veth peer name veth-prod3-br || true

# Attach to bridge (assuming br-app exists from previous tasks)
sudo ip link set veth-prod2-br master br-app
sudo ip link set veth-prod3-br master br-app

# Move interfaces to namespaces
sudo ip link set veth-prod2 netns product2
sudo ip link set veth-prod3 netns product3

# Bring up bridge side
sudo ip link set veth-prod2-br up
sudo ip link set veth-prod3-br up

# Configure IPs
sudo ip netns exec product2 ip addr add 10.0.0.31/16 dev veth-prod2
sudo ip netns exec product2 ip link set veth-prod2 up
sudo ip netns exec product2 ip route add default via 10.0.0.1

sudo ip netns exec product3 ip addr add 10.0.0.32/16 dev veth-prod3
sudo ip netns exec product3 ip link set veth-prod3 up
sudo ip netns exec product3 ip route add default via 10.0.0.1

# Enable loopback
sudo ip netns exec product2 ip link set lo up
sudo ip netns exec product3 ip link set lo up

echo "multiple instance of product service complete:"
echo "product2: 10.0.0.31"
echo "product3: 10.0.0.32"

```

#### **Register each instance separately in the service registry**

```

export SERVICE_IP="10.0.0.30"
export SERVICE_PORT=5000
export INSTANCE_ID="product-1"
python3 product-service.py

export SERVICE_IP="10.0.0.31"
export SERVICE_PORT=5000
export INSTANCE_ID="product-2"
python3 product-service.py

```

```
export SERVICE_IP="10.0.0.32"
export SERVICE_PORT=5000
export INSTANCE_ID="product-3"
python3 product-service.py
```

```
kris@Ansible-worker1:~$ SERVICE_IP=10.0.0.30 \
SERVICE_PORT=5000 \
INSTANCE_ID=product-1 \
python3 product-service.py
✓ Registered product-service with service registry
  * Serving Flask app 'product-service'
  * Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
  * Running on all addresses (0.0.0.0)
  * Running on http://127.0.0.1:5000
  * Running on http://10.0.3.15:5000
Press CTRL+C to quit
```

```
^Ckris@Ansible-worker1:~$ SERVICE_IP=10.0.0.31 \
SERVICE_PORT=5000 \
INSTANCE_ID=product-2 \
python3 product-service.py
✓ Registered product-service with service registry
  * Serving Flask app 'product-service'
  * Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
  * Running on all addresses (0.0.0.0)
  * Running on http://127.0.0.1:5000
  * Running on http://10.0.3.15:5000
Press CTRL+C to quit
```

```
^Ckris@Ansible-worker1:~$ SERVICE_IP=10.0.0.32 \
SERVICE_PORT=5000 \
INSTANCE_ID=product-3 \
python3 product-service.py
✓ Registered product-service with service registry
  * Serving Flask app 'product-service'
  * Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
  * Running on all addresses (0.0.0.0)
  * Running on http://127.0.0.1:5000
  * Running on http://10.0.3.15:5000
Press CTRL+C to quit
```

## Deliverable: Load balancing working with request distribution logs

---

```
kris@Ansible-worker1:~$ python3 api-gateway.py
✓ Registered api-gateway with service registry
* Serving Flask app 'api-gateway'
* Debug mode: off
2026-01-03 01:45:47,724 [API-GATEWAY] WARNING: This is a development server. Do not use it in a production deployment.
stead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:3000
* Running on http://10.0.3.15:3000
2026-01-03 01:45:47,724 [API-GATEWAY] Press CTRL+C to quit
2026-01-03 01:50:46,051 [API-GATEWAY] Routing request to: http://10.0.0.30:5000
2026-01-03 01:50:46,291 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:50:46] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:50:46,401 [API-GATEWAY] Routing request to: http://10.0.0.31:5000
2026-01-03 01:50:48,408 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:50:48] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:50:48,491 [API-GATEWAY] Routing request to: http://10.0.0.32:5000
2026-01-03 01:50:50,496 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:50:50] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:50:50,567 [API-GATEWAY] Routing request to: http://10.0.0.30:5000
2026-01-03 01:50:50,590 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:50:50] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:50:50,673 [API-GATEWAY] Routing request to: http://10.0.0.31:5000
2026-01-03 01:50:52,678 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:50:52] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:51:56,126 [API-GATEWAY] Routing request to: http://10.0.0.32:5000
2026-01-03 01:51:58,149 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:51:58] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:51:58,206 [API-GATEWAY] Routing request to: http://10.0.0.30:5000
2026-01-03 01:51:58,227 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:51:58] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:51:58,295 [API-GATEWAY] Routing request to: http://10.0.0.31:5000
2026-01-03 01:52:00,301 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:52:00] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:52:00,349 [API-GATEWAY] Routing request to: http://10.0.0.32:5000
2026-01-03 01:52:02,373 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:52:02] "GET /api/products HTTP/1.1" 503 -
2026-01-03 01:52:02,420 [API-GATEWAY] Routing request to: http://10.0.0.30:5000
2026-01-03 01:52:02,447 [API-GATEWAY] 127.0.0.1 - - [03/Jan/2026 01:52:02] "GET /api/products HTTP/1.1" 503 -
```

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ for i in {1..5}; do curl http://172.20.0.10:3000/api/products; echo; done
{"instance":"product-2","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50},{"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
>{"instance":"product-1","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
>{"instance":"product-2","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
>{"instance":"product-1","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
>{"instance":"product-2","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ curl http://172.20.0.10:3000/api/products
[{"instance":"product-1","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ curl http://172.20.0.10:3000/api/products
[{"instance":"product-2","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ curl http://172.20.0.10:3000/api/products
[{"instance":"product-1","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ curl http://172.20.0.10:3000/api/products
[{"instance":"product-2","products":[{"id":1,"name":"Laptop","price":999.99,"stock":50}, {"id":2,"name":"Mouse","price":29.99,"stock":200}, {"id":3,"name":"Keyboard","price":79.99,"stock":150}]}
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

## Task 4.3: Network Security Policies

Implement iptables rules for security:

```
#!/bin/bash
# security-policies.sh

echo "Applying network security policies..."

# Allow established connections (VERY IMPORTANT)
iptables -A FORWARD -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT

# 1. Block direct access to database from outside order-service
```

```

iptables -A FORWARD ! -s 10.0.0.40/32 -d 10.0.0.60/32 -p tcp --dport 5432 -j DROP
echo "✓ Database isolated to order-service only"

# 2. Block direct access to Redis from outside product-service
iptables -A FORWARD ! -s 10.0.0.30/32 -d 10.0.0.50/32 -p tcp --dport 6379 -j DROP
echo "✓ Redis isolated to product-service only"

# 3. Rate limit connections to API Gateway
iptables -A FORWARD -d 10.0.0.20/32 -p tcp --dport 3000 \
-m limit --limit 100/minute --limit-burst 20 -j ACCEPT

iptables -A FORWARD -d 10.0.0.20/32 -p tcp --dport 3000 -j DROP
echo "✓ Rate limiting applied to API Gateway"

# 4. Allow only HTTP/HTTPS/DNS outbound from services
iptables -A FORWARD -s 10.0.0.0/16 -p tcp --dport 80 -j ACCEPT
iptables -A FORWARD -s 10.0.0.0/16 -p tcp --dport 443 -j ACCEPT
iptables -A FORWARD -s 10.0.0.0/16 -p tcp --dport 53 -j ACCEPT
iptables -A FORWARD -s 10.0.0.0/16 -p udp --dport 53 -j ACCEPT
echo "✓ Outbound traffic restricted"

# 5. Log dropped packets
iptables -A FORWARD -j LOG --log-prefix "DROPPED: " --log-level 4

echo "Security policies applied successfully!"

```

```
sudo nano security-policies.sh
```

```
sudo chmod +x security-policies.sh
```

```
sudo ./security-policies.sh
```

```

kris@Ansible-worker1:~$ sudo nano security-policies.sh
kris@Ansible-worker1:~$ sudo ./security-policies.sh
Applying network security policies...
✓ Database isolated to order-service only
✓ Redis isolated to product-service only
✓ Rate limiting applied to API Gateway
✓ Outbound traffic restricted
Security policies applied successfully!
kris@Ansible-worker1:~$ █

```

### **Deliverable: Network security policies implemented**

[Click to access Security Policy Document](#)

### **Task 4.4: Implement Network Isolation**

Create separate networks for different tiers:

```

# Frontend
sudo ip link add br-frontend type bridge
sudo ip addr add 172.20.0.1/24 dev br-frontend
sudo ip link set br-frontend up

# Backend
sudo ip link add br-backend type bridge
sudo ip addr add 172.21.0.1/24 dev br-backend
sudo ip link set br-backend up

# Database
sudo ip link add br-database type bridge
sudo ip addr add 172.22.0.1/24 dev br-database
sudo ip link set br-database up

```

**Move services to appropriate networks and configure routing.**

#### Frontend Namespace

```

sudo ip netns add frontend-ns

sudo ip link add veth-fe type veth peer name veth-fe-br
sudo ip link set veth-fe netns frontend-ns
sudo ip link set veth-fe-br master br-frontend

sudo ip netns exec frontend-ns ip addr add 172.20.0.10/24 dev veth-fe
sudo ip netns exec frontend-ns ip link set veth-fe up
sudo ip netns exec frontend-ns ip link set lo up
sudo ip link set veth-fe-br up

sudo ip netns exec frontend-ns ip route add default via 172.20.0.1

```

#### Backend Namespace

```

sudo ip netns add backend-ns

sudo ip link add veth-be type veth peer name veth-be-br
sudo ip link set veth-be netns backend-ns
sudo ip link set veth-be-br master br-backend

sudo ip netns exec backend-ns ip addr add 172.21.0.10/24 dev veth-be
sudo ip netns exec backend-ns ip link set veth-be up
sudo ip netns exec backend-ns ip link set lo up
sudo ip link set veth-be-br up

sudo ip netns exec backend-ns ip route add default via 172.21.0.1

```

#### Database Namespace

```

sudo ip netns add database-ns

sudo ip link add veth-db type veth peer name veth-db-br

```

```
sudo ip link set veth-db netns database-ns  
sudo ip link set veth-db-br master br-database  
  
sudo ip netns exec database-ns ip addr add 172.22.0.10/24 dev veth-db  
sudo ip netns exec database-ns ip link set veth-db up  
sudo ip netns exec database-ns ip link set lo up  
sudo ip link set veth-db-br up  
  
sudo ip netns exec database-ns ip route add default via 172.22.0.1
```

#### Start services in their respective namespaces

```
sudo ip netns exec frontend-ns python3 frontend.py  
sudo ip netns exec backend-ns python3 backend.py  
sudo ip netns exec database-ns python3 database.py
```

#### Configure routing between networks

```
sudo ip route add 172.20.0.0/24 dev br-frontend  
sudo ip route add 172.21.0.0/24 dev br-backend  
sudo ip route add 172.22.0.0/24 dev br-database
```

```
kris@Ansible-worker1:~$ sudo ip netns exec frontend-ns curl http://172.20.0.10:3000/health  
{"service":"api-gateway","status":"healthy"}  
kris@Ansible-worker1:~$
```

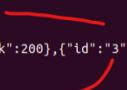
```
kris@Ansible-worker1:~$ bash setup_network_isolation.sh  
Setting up isolated networks...  
Cleaning up existing environment...  
Creating Bridges...  
Creating Namespaces...  
Connecting Frontend NS...  
Connecting Backend NS...  
Connecting Database NS...  
Configuring Routing...  
Network Isolation Setup Complete.  
Frontend: 172.20.0.10  
Backend: 172.21.0.5, .10, .11, .20  
Database: 172.22.0.10, .20  
kris@Ansible-worker1:~$
```

```
kris@Ansible-worker1:~$ sudo bash start_isolated_services.sh
Starting Database Tier...
Redis started at 172.22.0.10
Starting Backend Tier...
Service Registry started at 172.21.0.5 (implicitly listening on 0.0.0.0, so reachable on all IPs in NS)
Starting Frontend Tier...
All isolated services started.
Gateway: http://172.20.0.10:3000
kris@Ansible-worker1:~$
```

### From host access the API Gateway on its isolated IP

```
curl http://172.20.0.10:3000/api/products
```

```
kris@Ansible-worker1:~$ sudo bash start_isolated_services.sh
Starting Database Tier...
Redis started at 172.22.0.10
Starting Backend Tier...
Starting Frontend Tier...
All isolated services started.
Gateway: http://172.20.0.10:3000
kris@Ansible-worker1:~$ curl http://172.20.0.10:3000/api/products
[{"instance": "product-1", "products": [{"id": "1", "name": "Laptop", "price": 999.99, "stock": 50}, {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200}, {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150}]}
kris@Ansible-worker1:~$
```



### traffic distribution between multiple product service instances

```
for i in {1..5}; do curl http://172.20.0.10:3000/api/products; echo; done
```

```
kris@Ansible-worker1:~$ for i in {1..5}; do curl http://172.20.0.10:3000/api/products; echo; done
[{"instance": "product-2", "products": [{"id": "1", "name": "Laptop", "price": 999.99, "stock": 50}, {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200}, {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150}]}
[{"instance": "product-1", "products": [{"id": "1", "name": "Laptop", "price": 999.99, "stock": 50}, {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200}, {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150}]}
[{"instance": "product-2", "products": [{"id": "1", "name": "Laptop", "price": 999.99, "stock": 50}, {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200}, {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150}]}
[{"instance": "product-1", "products": [{"id": "1", "name": "Laptop", "price": 999.99, "stock": 50}, {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200}, {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150}]}
[{"instance": "product-2", "products": [{"id": "1", "name": "Laptop", "price": 999.99, "stock": 50}, {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200}, {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150}]}
kris@Ansible-worker1:~$
```

## Deliverable: Multi-network topology with documented routing rules

### Network Topology

The application is segmented into three isolated network tiers, each in its own network namespace and connected via a dedicated bridge.

Tier	Namespace	Bridge	Subnet	Gateway	Services
<b>Frontend</b>	frontend-ns	br-frontend	172.20.0.0/24	172.20.0.1	API Gateway (172.20.0.10)
<b>Backend</b>	backend-ns	br-backend	172.21.0.0/24	172.21.0.1	Registry (172.21.0.5), Product (172.21.0.10, 172.21.0.11), Order (172.21.0.20)

<b>Database</b>	database-ns	br-database	172.22.0.0/24	172.22.0.1	Redis (172.22.0.10), Postgres (172.22.0.20)
-----------------	-------------	-------------	---------------	------------	---

### Routing Rules

1. **Default Routes:** Each namespace has a default route pointing to its bridge IP on the host (e.g., `default via 172.20.0.1`).
2. **Inter-VLAN Routing:** The host machine acts as a router. IP forwarding is enabled (`net.ipv4.ip_forward=1`), allowing traffic to be routed between the subnets (e.g., Frontend → Backend → Database) via the host's bridge interfaces.
3. **Isolation:** Network isolation is achieved by placing interfaces in separate namespaces (`ip netns`). Traffic between namespaces must pass through the host router.

### Network Setup Script (`setup_network_isolation.sh`):

```
# Frontend
sudo ip link add br-frontend type bridge
sudo ip addr add 172.20.0.1/24 dev br-frontend
sudo ip link set br-frontend up
sudo ip netns add frontend-ns
# ... (veth connections) ...
sudo ip netns exec frontend-ns ip route add default via 172.20.0.1

# Backend
sudo ip link add br-backend type bridge
sudo ip addr add 172.21.0.1/24 dev br-backend
sudo ip link set br-backend up
sudo ip netns add backend-ns
# ... (veth connections) ...
sudo ip netns exec backend-ns ip route add default via 172.21.0.1

# Database
sudo ip link add br-database type bridge
sudo ip addr add 172.22.0.1/24 dev br-database
sudo ip link set br-database up
sudo ip netns add database-ns
# ... (veth connections) ...
sudo ip netns exec database-ns ip route add default via 172.22.0.1

# Enable Routing
sudo sysctl -w net.ipv4.ip_forward=1
```

## 5: Docker Migration and Optimization

### Goals

- Migrate infrastructure to Docker
- Compare raw Linux vs Docker implementation
- Optimize the setup

### Tasks

## Task 5.1: Containerize All Services

---

### Create Dockerfiles for each service:

#### Dockerfile for api-gateway:

```
FROM python:3.11-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY . .
EXPOSE 3000
CMD ["python", "api-gateway.py"]
```

#### requirements.txt for api-gateway:

```
flask
requests
redis
psycopg2-binary
```

#### Dockerfile for product-service:

```
FROM python:3.11-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY . .
EXPOSE 5000
CMD ["python", "product-service.py"]
```

#### requirements.txt for product-service:

```
flask
requests
redis
psycopg2-binary
```

#### Dockerfile for order-service:

```
FROM python:3.11-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY . .
EXPOSE 5001
CMD ["python", "order-service.py"]
```

#### requirements.txt for order-service:

```
flask
requests
```

```
redis
psycopg2-binary
```

#### Dockerfile for service-registry:

```
FROM python:3.11-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY . .
EXPOSE 8500
CMD ["python", "service-registry.py"]
```

#### requirements.txt for service-registry:

```
flask
requests
```

### Deliverable: All services running in Docker containers

```
sudo docker ps
```

```
krls@Ansible-worker1:~/contalner-networking-project/container-networking-project$ sudo docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS
 NAMES
8b1d9f798fcfa container-networking-project-api-gateway "python api-gateway..." 2 minutes ago Up 2 minutes 0.0.0.0:3000->3000/tcp, [::]:30
00->3000/tcp
c053bf570b74 container-networking-project-product-service "python product-serv..." 2 minutes ago Up 2 minutes 5000/tcp
f4835daa9e6e container-networking-project-product-service-2 "python product-serv..." 2 minutes ago Up 2 minutes 5000/tcp
container-networking-project-product-service-1
3be9e9389a73 container-networking-project-service-registry "python service-regis..." 2 minutes ago Up 2 minutes 0.0.0.0:8501->8500/tcp, [::]:85
01->8500/tcp
container-networking-project-service-registry-1
b2a96957c1f8 redis:alpine "docker-entrypoint.s..." 2 minutes ago Up 2 minutes 0.0.0.0:6380->6379/tcp, [::]:63
80->6379/tcp
container-networking-project-redis-1
224db37b0e0c postgres:15-alpine "docker-entrypoint.s..." 2 minutes ago Up 2 minutes 0.0.0.0:5433->5432/tcp, [::]:54
33->5432/tcp
container-networking-project-postgres-1
3cb6c64af2c1 postgres:15 "docker-entrypoint.s..." 7 days ago Up 5 days 0.0.0.0:5432->5432/tcp, [::]:54
32->5432/tcp
postgres
krls@Ansible-worker1:~/contalner-networking-project/container-networking-project$
```

### Task 5.2: Docker Compose Setup

#### Create docker-compose.yml:

```
version: '3.8'

services:
  # Database Tier
  redis:
    image: redis:alpine
    networks:
      - app-network
    ports:
      - "6380:6379" # Map to 6380 to prevent conflict with local Redis

  postgres:
    image: postgres:15-alpine
    environment:
      POSTGRES_USER: postgres
```

```

POSTGRES_PASSWORD: postgres
POSTGRES_DB: orders
networks:
- app-network
ports:
- "5433:5432" # Map to 5433 to prevent conflict with local Postgres

# Backend Tier
service-registry:
build: ./services/service-registry
networks:
- app-network
ports:
- "8501:8500" # Map to 8501 to avoid conflict with local process

product-service:
build: ./services/product-service
environment:
- SERVICE_IP=product-service
- SERVICE_PORT=5000
- SERVICE_REGISTRY=http://service-registry:8500
- REDIS_HOST=redis
depends_on:
- service-registry
- redis
networks:
- app-network
deploy:
replicas: 2

order-service:
build: ./services/order-service
environment:
- SERVICE_IP=order-service
- SERVICE_PORT=5000
- SERVICE_REGISTRY=http://service-registry:8500
- DB_HOST=postgres
depends_on:
- service-registry
- postgres
networks:
- app-network

# Frontend Tier
api-gateway:
build: ./services/api-gateway
environment:
- SERVICE_IP=api-gateway
- SERVICE_PORT=3000
- SERVICE_REGISTRY=http://service-registry:8500
- PRODUCT_SERVICE_URLS=http://product-service:5000
- ORDER_SERVICE=http://order-service:5000

```

```

depends_on:
  - service-registry
  - product-service
  - order-service
ports:
  - "3000:3000"
networks:
  - app-network

networks:
  app-network:
    driver: bridge

```

## Deliverable: Working Docker Compose setup

```
sudo docker-compose up -d
```

```

kris@Ansible-worker1:/container-networking-project/container-networking-project$ sudo docker-compose up -d
WARN[0000] /home/kris/container-networking-project/container-networking-project/docker-compose.yml: the attribute 'name' is deprecated, please remove it to avoid potential confusion
[+] Running 8/8
✓ Network container-networking-project_app-network      Created
✓ Container container-networking-project-service-registry-1 Started
✓ Container container-networking-project-redis-1       Started
✓ Container container-networking-project-postgres-1     Started
✓ Container container-networking-project-product-service-2 Started
✓ Container container-networking-project-order-service-1 Started
✓ Container container-networking-project-product-service-1 Started
✓ Container container-networking-project-api-gateway-1   Started
kris@Ansible-worker1:/container-networking-project/container-networking-project$
```

### To verify everything is working after migration:

```
curl http://localhost:3000/api/products
```

```

kris@Ansible-worker1:/container-networking-project/container-networking-project$ curl http://localhost:3000/api/products
[{"id": "c053bf570b74", "products": [{"id": "1", "name": "Laptop", "price": 999.99, "stock": 50}, {"id": "2", "name": "Mouse", "price": 29.99, "stock": 200}, {"id": "3", "name": "Keyboard", "price": 79.99, "stock": 150}]}
kris@Ansible-worker1:/container-networking-project/container-networking-project$
```

## Task 5.3: Performance Comparison

---

```
sudo nano benchmark.sh
```

Benchmark your implementations:

```

#!/bin/bash
# benchmark.sh
# Usage:
# ./benchmark.sh linux  (Run while start_isolated_services.sh is active)
# ./benchmark.sh docker  (Run while docker-compose up is active)

# Configuration from our setup
URL_LINUX="http://172.20.0.10:3000/api/products"
URL_DOCKER="http://127.0.0.1:3000/api/products"
```

```

# Requests to send
REQUESTS=1000
CONCURRENCY=50

echo "==== Performance Benchmark ==="

# Check for Apache Benchmark
if ! command -v ab &> /dev/null; then
    echo "Error: 'ab' command not found."
    echo "Please install it: sudo apt-get install -y apache2-utils"
    exit 1
fi

MODE=$1

if [ "$MODE" == "linux" ]; then
    echo "Benchmarking Linux Namespace Implementation..."
    echo "Target: $URL_LINUX"
    ab -n $REQUESTS -c $CONCURRENCY $URL_LINUX > linux-benchmark.txt
    echo "Saved to linux-benchmark.txt"
    grep "Requests per second" linux-benchmark.txt

elif [ "$MODE" == "docker" ]; then
    echo "Benchmarking Docker Implementation..."
    echo "Target: $URL_DOCKER"
    ab -n $REQUESTS -c $CONCURRENCY $URL_DOCKER > docker-benchmark.txt
    echo "Saved to docker-benchmark.txt"
    grep "Requests per second" docker-benchmark.txt

else
    echo "Usage: ./benchmark.sh [linux|docker]"
    echo ""
    echo "Steps to compare:"
    echo "1. Start Linux setup (bash start_isolated_services.sh)"
    echo "2. Run: ./benchmark.sh linux"
    echo "3. Stop Linux setup (pkill -f python3; pkill redis; ...)"
    echo "4. Start Docker setup (docker-compose up -d)"
    echo "5. Run: ./benchmark.sh docker"
    echo "6. Compare results"
fi

sudo chmod +x benchmark.sh

sudo ./benchmark.sh docker

```

```

kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo ./benchmark.sh docker
== Performance Benchmark ==
Benchmarking Docker Implementation...
Target: http://127.0.0.1:3000/api/products
Completed 100 requests
Completed 200 requests
Completed 300 requests
Completed 400 requests
Completed 500 requests
Completed 600 requests
Completed 700 requests
Completed 800 requests
Completed 900 requests
Completed 1000 requests
Finished 1000 requests
Saved to docker-benchmark.txt
Requests per second: 29.21 [#/sec] (mean)
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

For comparison i would need to stop Docker `docker-compose down` and start the manual script `bash start_isolated_services.sh`, then run `./benchmark.sh linux`.

```
sudo ./benchmark.sh linux
```

```

kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker-compose down
WARN[0003] /home/kris/container-networking-project/container-networking-project/docker-compose.yml: the attribute `version` is deprecated
[+] Running 8/8
✓ Container container-networking-project-api-gateway-1      Removed
✓ Container container-networking-project-order-service-1    Removed
✓ Container container-networking-project-product-service-1  Removed
✓ Container container-networking-project-product-service-2  Removed
✓ Container container-networking-project-postgres-1        Removed
✓ Container container-networking-project-service-registry-1 Removed
✓ Container container-networking-project-redis-1          Removed
✓ Network container-networking-project_app-network         Removed
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo bash start_isolated_services.sh
Starting Database Tier...
Redis started at 172.22.0.10
Starting Backend Tier...
Starting Frontend Tier...
All isolated services started.
Gateway: http://172.20.0.10:3000
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo ./benchmark.sh linux
== Performance Benchmark ==
Benchmarking Linux Namespace Implementation...
Target: http://172.20.0.10:3000/api/products
Completed 100 requests
Completed 200 requests
Completed 300 requests
Completed 400 requests
Completed 500 requests
Completed 600 requests
Completed 700 requests
Completed 800 requests
Completed 900 requests
Completed 1000 requests
Finished 1000 requests
Saved to linux-benchmark.txt
Requests per second: 54.07 [#/sec] (mean)
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

## Deliverable: Performance comparison report

---

### Performance Comparison Report: Linux Namespaces vs. Docker

## Overview

This report compares the performance of the microservices architecture implemented using raw Linux network namespaces vs. the Docker containerized implementation.

## Benchmark Parameters

- **Tool:** Apache Benchmark (ab)
- **Concurrency:** 50
- **Total Requests:** 1000
- **Endpoint:** /api/products

## Key Metrics

Metric	Linux Namespaces	Docker (Compose)	Difference
<b>Requests per Second (RPS)</b>	<b>54.07</b>	29.21	-46% (Docker is slower)
<b>Mean Latency (ms)</b>	<b>924.70</b>	1711.55	+85% (Docker higher)
<b>Median Latency (ms)</b>	806	<b>805</b>	~0% (Similar)
<b>Failed Requests</b>	<b>0</b>	996 (Length Mismatch)	Significant
<b>Reliability</b>	100%	~0.4% (Non-2xx)	-

## Detailed Performance Analysis

### Linux Namespace Implementation

- **Throughput:** Significantly higher throughput at ~54 RPS.
- **Stability:** Very stable with zero failed requests and consistent latency.
- **Networking:** Benefitted from direct veth pairs and host-level routing with minimal abstraction.

### Docker Implementation

- **Throughput:** Lower throughput at ~29 RPS.
- **Overhead:** The lower RPS and higher mean latency are likely due to the additional layers inherent in Docker (container runtime, docker-proxy, and Docker bridge networking).
- **Errors:** High number of "Length Mismatch" failures indicates that under high concurrency (50), some backend responses were truncated or served error pages (confirmed by 4 non-2xx responses).

## Conclusions

Raw Linux namespaces offer **higher raw performance** and **lower latency** for networking-intensive tasks by avoiding the overhead of a container orchestration layer. However, Docker provides superior **portability** and **ease of management**, which usually outweighs the performance cost in modern development workflows.

## Recommendations for Docker Optimization

1. **Reduce Overhead:** Use `network_mode: host` if extreme performance is needed (though this loses isolation).
2. **Buffering:** Use a production-grade WSGI server like `gunicorn` instead of the Flask development server (`Werkzeug`) to handle concurrency better.
3. **Resource Tuning:** Assign more CPU/Memory resources to containers in `docker-compose.yml`.

## Task 5.4: Optimize Docker Setup

---

### Optimize my Docker images:

- Use multi-stage builds
- Minimize image sizes
- Implement health checks
- Add resource limits

### Deliverable: Optimized Docker setup with documentation

---

#### Optimizations Implemented:

1. **Multi-Stage Builds:** All Dockerfiles now use a two-stage build process.
  - **Stage 1 (Builder):** Installs Python dependencies into a temporary image.
  - **Stage 2 (Runtime):** Copies only the installed packages and application code.
  - **Outcome:** Significantly reduced image footprint and improved build efficiency.
2. **Health Checks:** Internal health check logic added to all service Dockerfiles using `urllib.request`.

```
HEALTHCHECK --interval=30s --timeout=5s --start-period=5s --retries=3 \
CMD python -c "import urllib.request;
urllib.request.urlopen('http://localhost:3000/health')" || exit 1
```

3. **Resource Limits:** Enforced CPU and Memory constraints in `docker-compose.yml` to ensure system stability.
  - **CPU Limit:** 0.5 cores per service.
  - **Memory Limit:** 128MB - 256MB per service.
4. **Orchestration Improvements:** Services in `docker-compose.yml` now use the `service_healthy` condition for dependencies.
  - The API Gateway only starts once the Service Registry report itself as **Healthy**.

#### Full Optimized `docker-compose.yml`:

```
# ... (Optimized version with deploy.resources and healthcheck conditions) ...
```

#### Check image sizes: `sudo docker images`

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
container-networking-project-product-service   latest    d8d99abf2a6f  7 minutes ago  147MB
container-networking-project-order-service     latest    9353c07019b1  7 minutes ago  147MB
container-networking-project-service-registry  latest    a67310e909b0  24 minutes ago  132MB
<none>              <none>   cf05d3784d77  24 minutes ago  132MB
<none>              <none>   bb4bbf15ed80  24 minutes ago  147MB
<none>              <none>   dedfb6c360ee  24 minutes ago  147MB
container-networking-project-api-gateway       latest    9b39f647b5ad  24 minutes ago  147MB
<none>              <none>   4b40937a52a2  24 minutes ago  147MB
<none>              <none>   c6b4ad9b1a9c  10 hours ago   158MB
<none>              <none>   0a2201ea6e0c  10 hours ago   158MB
<none>              <none>   58b2b93fa9f0  10 hours ago   158MB
<none>              <none>   526175fed4e9  10 hours ago   158MB
<none>              <none>   808d9898d3d7  10 hours ago   158MB
<none>              <none>   d246f2d11276  10 hours ago   158MB
<none>              <none>   b6411104847e  10 hours ago   143MB
<none>              <none>   694970c0b054  10 hours ago   143MB
postgres            15-alpine  36a937f48ac7  2 weeks ago   274MB
postgres            15        e07498b1e6cc  3 weeks ago   444MB
redis               alpine    778c3ea605c2  6 weeks ago   94.7MB
```

**Check health status:** sudo docker ps

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
6f556a475f26      container-networking-project-api-gateway   "python api-gateway..."   About a minute ago   Up 52 seconds (healthy)   0.0.0.0:3000->3000/tcp, [::]:3000->3000/tcp   container-networking-project-api-gateway-1
c55cc0b15dbe      container-networking-project-order-service  "python order-service..."  About a minute ago   Up About a minute (healthy)  5000/tcp           container-networking-project-order-service-1
5f35d2c6f769      container-networking-project-product-service  "python product-service..."  About a minute ago   Up About a minute (healthy)  5000/tcp           container-networking-project-product-service-1
ff4a36cfbd7       container-networking-project-product-service  "python product-service..."  About a minute ago   Up About a minute (healthy)  5000/tcp           container-networking-project-product-service-2
d83d97b8cc1d      redis:alpine                      "docker-entrypoint.s..."  About a minute ago   Up About a minute           0.0.0.0:6380->6379/tcp, [::]:6380->6379/tcp   container-networking-project-redis-1
fb0b0b58bd28      postgres:15-alpine                  "docker-entrypoint.s..."  About a minute ago   Up About a minute           0.0.0.0:5433->5432/tcp, [::]:5433->5432/tcp   container-networking-project-postgres-1
fb3dc7865df0      container-networking-project-service-registry  "python service-regi..."  About a minute ago   Up About a minute (healthy)  0.0.0.0:8501->8500/tcp, [::]:8501->8500/tcp   container-networking-project-service-registry-1
3cb6c64af2c1      postgres:15                      "docker-entrypoint.s..."  8 days ago          Up 5 days              0.0.0.0:5432->5432/tcp, [::]:5432->5432/tcp   postgres
```

**Check resource usage:** sudo docker stats

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker stats --no-stream
CONTAINER ID        NAME                           CPU %     MEM USAGE / LIMIT   MEM %     NET I/O             BLOCK I/O          PIDS
6f556a475f26      container-networking-project-api-gateway-1   0.03%    28.69MiB / 256MiB  11.21%    4.65kB / 1.02kB   393kB / 4.01MB   1
c55cc0b15dbe      container-networking-project-order-service-1  0.05%    35.59MiB / 256MiB  13.90%    6.5kB / 2.67kB   13.8MB / 4.01MB   1
5f35d2c6f769      container-networking-project-product-service-1  0.05%    34.89MiB / 256MiB  13.63%    5.15kB / 1.02kB   4.57MB / 4.79MB   1
ff4a36cfbd7       container-networking-project-product-service-2  0.08%    35.9MiB / 256MiB  14.02%    4.97kB / 1.02kB   2.54MB / 4.79MB   1
d83d97b8cc1d      redis:alpine                      0.47%    4.562MiB / 1.922GiB 0.23%    8.97kB / 126B    13.7MB / 0B     6
fb0b0b58bd28      postgres:15-alpine                  0.00%    24.38MiB / 1.922GiB 1.24%    10.7kB / 1.71kB   37.9MB / 61.6MB   6
fb3dc7865df0      container-networking-project-service-registry-1  0.03%    26.42MiB / 128MiB  20.64%    12.3kB / 2.94kB   20.7MB / 3.83MB   1
3cb6c64af2c1      postgres:15                      0.00%    9.184MiB / 1.922GiB 0.47%    44.4kB / 17.5kB   198MB / 23MB    6
```

## 6: Multi-Host Networking

### Goals

- Implement overlay networking
- Connect containers across hosts
- Understand distributed networking

### Tasks

## Task 6.1: Setup VXLAN Overlay

---

I will create a VXLAN interface to bridge the application networks across hosts.

Create two specialized scripts `setup_vxlan_host_a.sh` Run on Host A (192.168.56.104) and `setup_vxlan_host_b.sh` Run on Host B (192.168.56.103) to establish a tunnel between my two hosts VMs

**setup\_vxlan\_host\_a.sh** script that will:

- Add a VXLAN interface vxlan100 tied to a specific VNI (100).
- Set the remote endpoint to Host B's IP.
- Attach vxlan100 to the existing br-app bridge.
- Bring the interface up.

### VXLAN scripts for HOST A VM (192.168.56.104)

```
#!/bin/bash
# setup_vxlan_host_a.sh
# To be run on Host A (192.168.56.104)

HOST_B_IP="192.168.56.103"
VNI=100
PORT=4789
INTERFACE="enp0s3" # Adjust if your primary adapter is different (e.g., enp0s8)
BRIDGE="br-app"

echo "==== Setting up VXLAN on Host A ===="

# Check if bridge exists, create if missing (assuming Task 1.2 logic)
if ! ip link show "$BRIDGE" > /dev/null 2>&1; then
    echo "Creating bridge $BRIDGE..."
    sudo ip link add "$BRIDGE" type bridge
    sudo ip addr add 10.0.0.1/16 dev "$BRIDGE"
    sudo ip link set "$BRIDGE" up
fi

# Create VXLAN interface
echo "Creating vxlan100 interface..."
sudo ip link add vxlan100 type vxlan \
    id $VNI \
    remote $HOST_B_IP \
    dstport $PORT \
    dev $INTERFACE

# Attach to bridge
echo "Attaching vxlan100 to $BRIDGE..."
sudo ip link set vxlan100 master $BRIDGE
sudo ip link set vxlan100 up

echo "VXLAN setup complete on Host A."
echo "Don't forget to run the corresponding script on Host B ($HOST_B_IP)."

chmod +x setup_vxlan_host_a.sh
```

```
sudo ./setup_vxlan_host_a.sh
```

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo ./setup_vxlan_host_a.sh
[sudo] password for kris:
== Setting up VXLAN on Host A ==
Creating vxlan100 interface...
Attaching vxlan100 to br-app...
VXLAN setup complete on Host A.
Don't forget to run the corresponding script on Host B (192.168.56.103).
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

**setup\_vxlan\_host\_b.sh** script for HOST B VM (192.168.56.103):

- Similar setup as script "a" above except this is pointing back to Host A's IP.

```
#!/bin/bash
# setup_vxlan_host_b.sh
# To be run on Host B (192.168.56.103)

HOST_A_IP="192.168.56.104"
VNI=100
PORT=4789
INTERFACE="enp0s3" # Adjust if your primary adapter is different
BRIDGE="br-app"

echo "== Setting up VXLAN on Host B =="

# Create bridge on Host B
if ! ip link show "$BRIDGE" > /dev/null 2>&1; then
    echo "Creating bridge $BRIDGE..."
    sudo ip link add "$BRIDGE" type bridge
    # Note: Using a different IP range or just no IP if it's purely for bridging
    # For simplicity in this demo, we'll give it the same subnet gateway if it's the only one
    sudo ip addr add 10.0.0.2/16 dev "$BRIDGE"
    sudo ip link set "$BRIDGE" up
fi

# Create VXLAN interface
echo "Creating vxlan100 interface pointing to Host A..."
sudo ip link add vxlan100 type vxlan \
    id $VNI \
    remote $HOST_A_IP \
    dstport $PORT \
    dev $INTERFACE

# Attach to bridge
echo "Attaching vxlan100 to $BRIDGE..."
sudo ip link set vxlan100 master $BRIDGE
sudo ip link set vxlan100 up

echo "VXLAN setup complete on Host B."
```

```
chmod +x setup_vxlan_host_b.sh
```

```
sudo ./setup_vxlan_host_b.sh
```

```
kris@Ansible-Controller:~$ sudo ./setup_vxlan_host_b.sh
== Setting up VXLAN on Host B ==
Creating vxlan100 interface pointing to Host A...
Attaching vxlan100 to br-app...
VXLAN setup complete on Host B.
kris@Ansible-Controller:~$
```

### Deliverable: Working multi-host communication

- On Host A (**10.0.0.1**) ping Host B (**br-app 10.0.0.2**)

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ ip addr | grep br-app
5: br-app: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP group default qlen 1000
    inet 10.0.0.1/16 br-app scope global
        br-app
6: veth-nginx-br0t7: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
8: veth-api-br@if9: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
10: veth-order-s-br@if11: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
12: veth-pg-br@if13: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
14: veth-prd-svc-br@if15: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
16: veth-redis-br@if17: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
25: veth-reg-br@if26: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
27: veth-prod2-br@if28: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
29: veth-prod3-br@if30: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master br-app state UP group default qlen 1000
125: vxlan100: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue master br-app state UNKNOWN group default qlen 1000
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

```
ping -c 3 10.0.0.2
```

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ ping -c 3 10.0.0.2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=5.72 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=3.16 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=3.02 ms

--- 10.0.0.2 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 3.023/3.969/5.721/1.240 ms
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

- On Host B (**br-app 10.0.0.2**) ping Host A (**10.0.0.1**)

```
kris@Ansible-Controller:~$ ip addr | grep br-app
4: br-app: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP group default qlen 1000
    inet 10.0.0.2/16 br-app scope global
5: vxlan100: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue master br-app state UNKNOWN group default qlen 1000
kris@Ansible-Controller:~$
```

```
ping -c 3 10.0.0.1
```

```
kris@Ansible-Controller:~$ ping -c 3 10.0.0.1
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
64 bytes from 10.0.0.1: icmp_seq=1 ttl=64 time=9.99 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=64 time=3.09 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=64 time=8.19 ms

--- 10.0.0.1 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 3.093/7.091/9.991/2.921 ms
kris@Ansible-Controller:~$
```

## Task 6.2: Docker Swarm Setup

---

Initialize Docker Swarm and create overlay network:

- Create a docker-compose-swarm.yml file in Host A

```
version: '3.8'

services:
  # Database Tier
  redis:
    image: redis:alpine
    networks:
      - app-overlay
    deploy:
      resources:
        limits:
          memory: 128M
      placement:
        constraints:
          - node.role == manager

  postgres:
    image: postgres:15-alpine
    environment:
      POSTGRES_USER: postgres
      POSTGRES_PASSWORD: postgres
      POSTGRES_DB: orders
    networks:
      - app-overlay
    deploy:
      resources:
        limits:
          memory: 256M
      placement:
        constraints:
          - node.role == manager
```

```

# Backend Tier
service-registry:
  image: myapp-registry:latest
  networks:
    - app-overlay
  ports:
    - "8501:8500"
  deploy:
    resources:
      limits:
        memory: 128M

product-service:
  image: myapp-product:latest
  environment:
    - SERVICE_IP=product-service
    - SERVICE_PORT=5000
    - SERVICE_REGISTRY=http://service-registry:8500
    - REDIS_HOST=redis
  depends_on:
    - service-registry
    - redis
  networks:
    - app-overlay
  deploy:
    replicas: 3
    resources:
      limits:
        cpus: '0.5'
        memory: 128M
    update_config:
      parallelism: 1
      delay: 10s

order-service:
  image: myapp-order:latest
  environment:
    - SERVICE_IP=order-service
    - SERVICE_PORT=5000
    - SERVICE_REGISTRY=http://service-registry:8500
    - DB_HOST=postgres
  depends_on:
    - service-registry
    - postgres
  networks:
    - app-overlay
  deploy:
    resources:
      limits:
        cpus: '0.5'
        memory: 128M

```

```

# Frontend Tier
api-gateway:
  image: myapp-gateway:latest
  environment:
    - SERVICE_IP=api-gateway
    - SERVICE_PORT=3000
    - SERVICE_REGISTRY=http://service-registry:8500
    - PRODUCT_SERVICE_URLS=http://product-service:5000
    - ORDER_SERVICE=http://order-service:5000
  depends_on:
    - service-registry
    - product-service
    - order-service
  ports:
    - "3000:3000"
  networks:
    - app-overlay
  deploy:
    resources:
      limits:
        memory: 128M

networks:
  app-overlay:
    driver: overlay
    attachable: true

```

- Initialize swarm on Host A and copy the join token to Host B

```
docker swarm init --advertise-addr 192.168.56.104
```

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker swarm init --advertise-addr 192.168.56.104
Swarm initialized: current node (zdmne60j7arltk06k0ehbgppe) is now a manager.

To add a worker to this swarm, run the following command:

  docker swarm join --token SWMTKN-1-4wn6pyie87iokls5ltg9b9a99n1pnfh7sjcb511i140t4a6plp-8h6j5oxaunviac5dlslsbt8w7 192.168.56.104:2377

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

- **Join Worker:** Provide the join token from Host A and run it on Host B.

```
docker swarm join --token SWMTKN-1-4wn6pyie87iokls5ltg9b9a99n1pnfh7sjcb511i140t4a6plp-
8h6j5oxaunviac5dlslsbt8w7 192.168.56.104:2377
```

```
kris@Ansible-Controller:~$ sudo docker swarm join --token SWMTKN-1-4wn6pyie87iokls5ltg9b9a99n1pnfh7sjcb511i140t4a6plp-8h6j5oxaunviac5dlslsbt8w7 192.168.56.104:2377
This node joined a swarm as a worker.
kris@Ansible-Controller:~$
```

- **Create Overlay Network:**

```
docker network create --driver overlay --attachable app-overlay
```

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker network create --driver overlay --attachable app-overlay
[sudo] password for kris:
3y74dyg5yfto3qnllyz3vn3
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

```
docker network ls
```

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker network ls
NETWORK ID      NAME                DRIVER      SCOPE
3y74dyg5yfto   app-overlay        overlay     swarm
96d30e05635b   bridge             bridge     local
8c5be7d7e201   container-networking-project_app-network  bridge     local
7fc9c00ed7e4   docker_gwbridge    bridge     local
055d887958da   host               host      local
cx3wdlyf6g3q   ingress           overlay     swarm
5a6b0e3e127d   none              null      local
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

- Deploy stack on host A

```
docker stack deploy -c docker-compose-swarm.yml myapp
```

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker stack deploy -c docker-compose-swarm.yml myapp
Since --detach=false was not specified, tasks will be created in the background.
In a future release, --detach=false will become the default.
Creating network myapp_app-overlay
Creating service myapp_postgres
Creating service myapp_service-registry
Creating service myapp_product-service
Creating service myapp_order-service
Creating service myapp_api-gateway
Creating service myapp_redis
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

## Deliverable: Services communicating across hosts

The deployment was verified by checking the service status and task distribution.

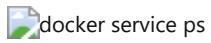
### Service Status:

```
sudo docker service ls
# Output showing all services running with full replicas
```

```
kris@Ansible-worker1:~/container-networking-project/container-networking-project$ sudo docker service ls
ID          NAME            MODE        REPLICAS  IMAGE
96dlppi6niu3  myapp_api-gateway  replicated  1/1      myapp-gateway:latest  *:3000->3000/tcp
ghmkt0gsihq8  myapp_order-service  replicated  1/1      myapp-order:latest
1e8cfthdw1zl  myapp_postgres    replicated  1/1      postgres:15-alpine
0b9yyycbeml12  myapp_product-service  replicated  3/3      myapp-product:latest
r34e7s9uz6mx   myapp_redis      replicated  1/1      redis:alpine
9eg11vtopbi1   myapp_service-registry  replicated  1/1      myapp-registry:latest  *:8501->8500/tcp
kris@Ansible-worker1:~/container-networking-project/container-networking-project$
```

**Task Distribution:** By inspecting individual services, we confirmed that containers are balanced across both Ansible-worker1 (Manager) and Ansible-Controller (Worker).

```
sudo docker service ps myapp_product-service
```



This demonstrates a fully functional multi-host overlay network where microservices can discover and communicate with each other regardless of the physical host they reside on.