

Technical Document: Container Networking Project

1. Architecture Document

System Overview

This project demonstrates a microservices architecture implemented using two distinct containerization approaches:

1. **Lower-level Linux Primitives:** Utilizing network namespaces, veth pairs, bridges, and iptables.
2. **Docker Containerization:** Utilizing Docker Compose, multi-stage builds, and custom bridge networks.

The application is an e-commerce platform consisting of an API Gateway, a Service Registry, a Product Service (with caching), an Order Service (with persistence), Redis, and PostgreSQL.

Component Descriptions

- **API Gateway (Python/Flask):** Entry point for external traffic. Performs round-robin load balancing to backend instances.
- **Service Registry (Python/Flask):** Simplified DNS-like registry for service discovery.
- **Product Service (Python/Flask):** Manages product data; uses Redis for caching to improve performance.
- **Order Service (Python/Flask):** Manages customer orders; persists data in PostgreSQL.
- **Redis Cache:** High-performance key-value store used for session and response caching.
- **PostgreSQL:** Relational database for persistent storage of orders.

Network Topology

The system is divided into three isolated security tiers:

- **Frontend (172.20.0.0/24):** Houses the API Gateway.
- **Backend (172.21.0.0/24):** Houses logic services (Product, Order, Registry).
- **Database (172.22.0.0/24):** Houses data stores (Redis, Postgres).

In the Linux implementation, each tier has a dedicated bridge (`br-frontend` , `br-backend` , `br-database`) and the host acts as the router. In the Docker implementation, all services reside on a unified `app-network` bridge with internal DNS discovery.

Data Flow Diagram



2. Implementation Guide

Step-by-Step Setup

Approach A: Linux Namespaces

1. **Configure Network:** Run `sudo ./setup_network_isolation.sh` to create bridges and namespaces.
2. **Start Services:** Run `sudo bash start_isolated_services.sh` to launch services inside namespaces.
3. **Verify:** Use `curl http://172.20.0.10:3000/api/products` from the host.

Approach B: Docker Compose

1. **Build & Launch:** Run `docker-compose up --build -d`.
2. **Verify Health:** Run `docker ps` to ensure status is "healthy".
3. **Test Endpoint:** Use `curl http://localhost:3000/api/products`.

Configuration Files

- **docker-compose.yml:** Defines service dependencies, resource limits, and environment variables.
- **Dockerfiles:** Use multi-stage builds (`python:3.11-slim`) to minimize image size and include built-in healthchecks.
- **setup_network_isolation.sh:** Shell script utilizing `ip netns` and `brctl` (or `ip link bridge`) for manual isolation.

Troubleshooting Guide

- **Connection Refused:** Ensure services are bound to `0.0.0.0` inside containers.
- **Registry Failures:** Check that the `SERVICE_REGISTRY` environment variable matches the internal network IP/hostname.
- **Healthcheck Unhealthy:** Verify the `/health` endpoint is reachable from within the container context.

3. Operations Manual

How to Start/Stop Services

- **Linux Setup:**
 - Start: `sudo bash start_isolated_services.sh`
 - Stop: `pkill -f python3` and `sudo ip netns del <name>`
- **Docker Setup:**
 - Start: `docker-compose up -d`
 - Stop: `docker-compose down`

Monitoring Procedures

- **Traffic Logs:** Run `./monitor-traffic.sh` to capture live pcap/text logs from the bridge.
- **Service Status:**
 - Linux: `sudo ip netns exec <ns> ss -lntp`
 - Docker: `docker ps`
- **Resource Usage:** `docker stats` for CPU/Memory consumption.

Backup and Recovery

- **Database Backup:** Use `docker exec postgres pg_dump -U postgres orders > backup.sql`.
- **Restoration:** `cat backup.sql | docker exec -i postgres psql -U postgres orders`.

Scaling Guidelines

- **Horizontal Scaling:** Use `docker-compose up --scale product-service=3 -d` to add more instances.
The API Gateway will automatically load balance via the service name alias.

4. Comparison Analysis

Linux Primitives vs Docker

Feature	Linux Namespaces	Docker Compose
Complexity	High (Shell scripts, manual routing)	Low (YAML description)
Overhead	Minimal (Direct kernel features)	Moderate (Runtime + Proxy)
Isolation	Strong networking isolation	Standard container isolation
Ease of Deployment	Low (Environment specific)	High (Portable)

Performance Metrics

- **Throughput (RPS):** Linux (54.07) > Docker (29.21)
- **Mean Latency:** Linux (924ms) < Docker (1711ms)
- **Observation:** Raw Linux namespaces are ~46% faster in high-concurrency scenarios due to lack of proxying overhead.

Pros and Cons

Linux Primitives:

- *Pros:* Maximum performance, educational value, deep control.
- *Cons:* Difficult to scale, brittle scripts, high maintenance.

Docker:

- *Pros:* Rapid scaling, dependency management, industry standard.
- *Cons:* Performance tax (networking hop), larger image sizes if not optimized.